An Experimental Evaluation of In-Memory Multi-Version Concurrency Control

Yingjun Wu, Joy Arulraj, Jiexi Lin, Ran Xian, Andrew Pavlo

MULTI-VERSION CONCURRENCY CONTROL

- Avoid read-write conflicts
- Support time-travel queries
- Enable snapshot isolation

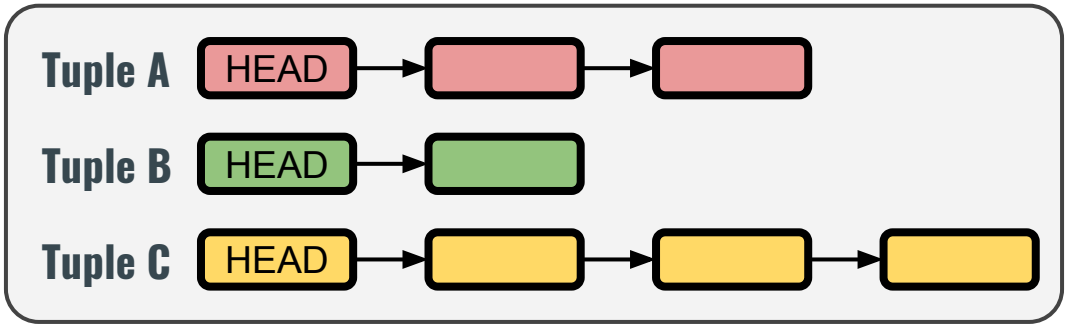

VERSION CHAINS

A BRIEF HISTORY OF MVCC

1979: FIRST MENTION

DAVID REED

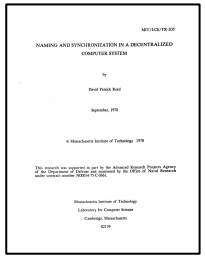

MIT/LCS/TR-205

NAMING AND SYNCHRONIZATION IN A DECENTRALIZED COMPUTER SYSTEM

by

David Patrick Reed

September, 1978

© Massachusetts Institute of Technology 1978

This research was supported in part by the Advanced Research Projects Agency of the Department of Defense and monitored by the Office of Naval Research under contract number N00014-75-C-0661.

Massachusetts Institute of Technology
Laboratory for Computer Science
Cambridge, Massachusetts
02139

Naming and synchronization in a decentralized computer system
Ph.D. Thesis, 1979

A BRIEF HISTORY OF MVCC

- 1979: FIRST MENTION
- 1981: FIRST IMPLEMENTATION

InterBase/Firebird

A BRIEF HISTORY OF MVCC

- 1979: **FIRST MENTION**
- 1981: **FIRST IMPLEMENTATION**
- 1984: Oracle
- 1985: Postgres

A BRIEF HISTORY OF MVCC

- 1979: **FIRST MENTION**
- 1981: **FIRST IMPLEMENTATION**
- 1984: Oracle
- 1985: Postgres
- 2001: MySQL-InnoDB

A BRIEF HISTORY OF MVCC

- 1979: **FIRST MENTION**
- 1981: **FIRST IMPLEMENTATION**
- 1984: Oracle
- 1985: Postgres
- 2001: MySQL-InnoDB
- 2010-2017: Hyrise, Hekaton, MemSQL, SAP HANA, NuoDB, HyPer...

A BRIEF HISTORY OF MVCC

Search for the best MVCC scheme for multi-core main-memory DBMSs

DESIGN DECISIONS OF MVCC

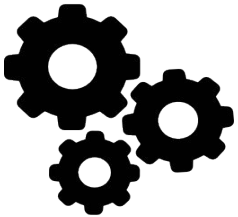

Concurrency Control Protocol

Version Storage

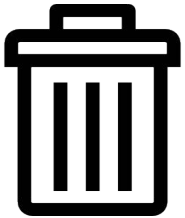

Garbage Collection

Index Management

CONCURRENCY CONTROL PROTOCOL

SCHEME	DBMS
Timestamp Ordering (MVTO)	N/A
Optimistic Concurrency Control (MVOCC)	   
Two-phase Locking (MV2PL)	   
Serialization Certifier	

CONCURRENCY CONTROL PROTOCOL

- **Approach #1: Timestamp Ordering (MVTO)**
  - The DBMS assigns transactions timestamps that determine serial order.

	TXN-ID	READ-TS	BEGIN-TS	END-TS
A_x	0	Tid	10	20
B_x	Tid	17	15	30
B_{x+1}	Tid	0	-	-

CONCURRENCY CONTROL PROTOCOL

- **Approach #3: Optimistic Concurrency Control (MVOCC)**
  - Transactions optimistically access physical versions before validating the read-set consistency.

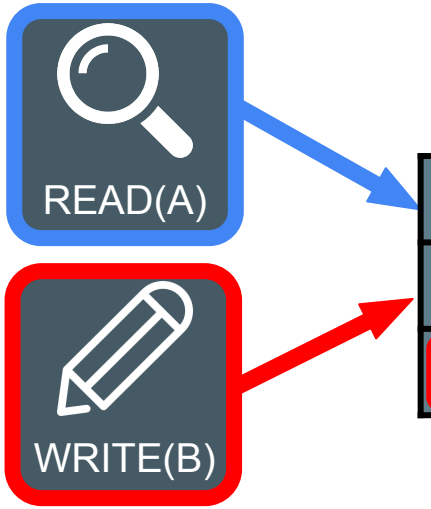

	TXN-ID	BEGIN-TS	END-TS
A_x	0	10	20
B_x	Tid	15	30
B_{x+1}	Tid	-	-

CONCURRENCY CONTROL PROTOCOL

- **Approach #3: Two-Phase Locking (MV2PL)**
  - Transactions acquire appropriate lock on physical version before they can read/write a logical tuple.

	TXN-ID	READ-CNT	BEGIN-TS	END-TS
A_x	0	2	10	20
B_x	Tid	0	15	30
B_{x+1}	Tid	0	-	-

CONCURRENCY CONTROL PROTOCOL

- **Approach #4: Serialization Certifier**
  - The DBMS maintains a serialization graph for detecting and removing "dangerous structures" formed by concurrent transactions.

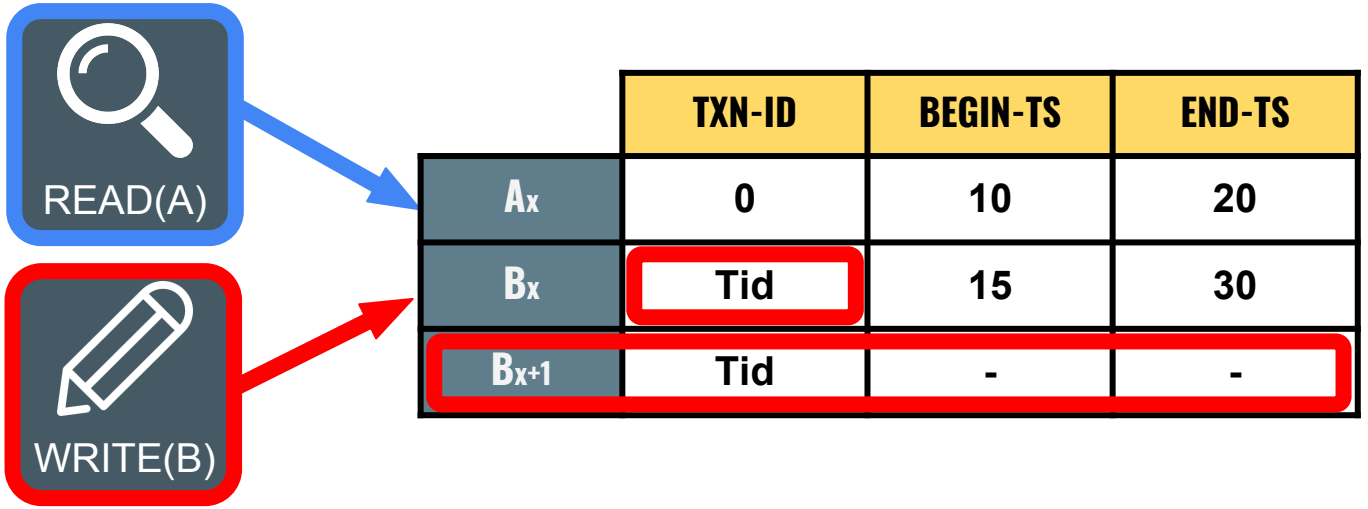

VERSION STORAGE

SCHEME	DBMS
Append-Only	    
Time-Travel	
Delta	  

VERSION STORAGE

- Approach #1: Append-Only Storage (Oldest-to-Newest)
  - New versions are appended to the same table space.

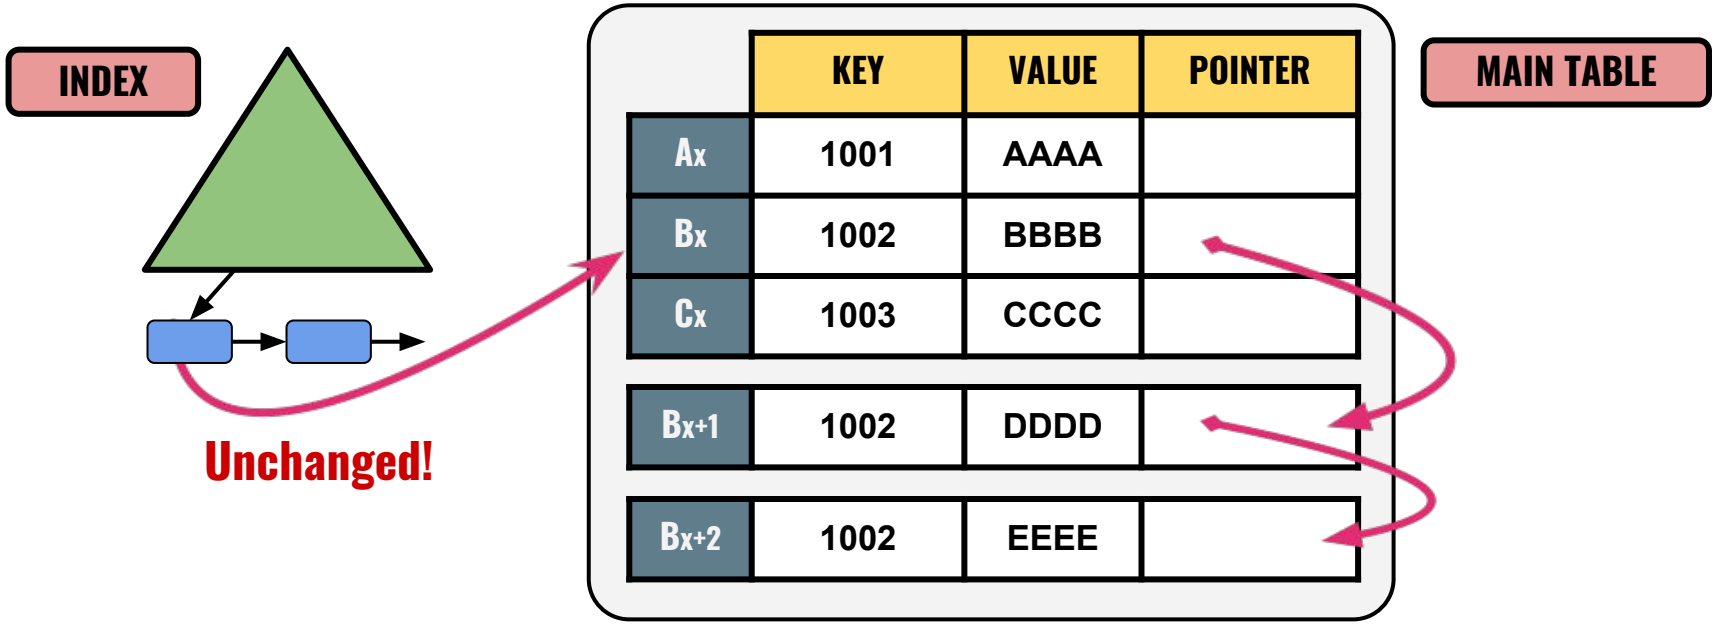

VERSION STORAGE

- Approach #1: Append-Only Storage (Newest-to-Oldest)
  - New versions are appended to the same table space.

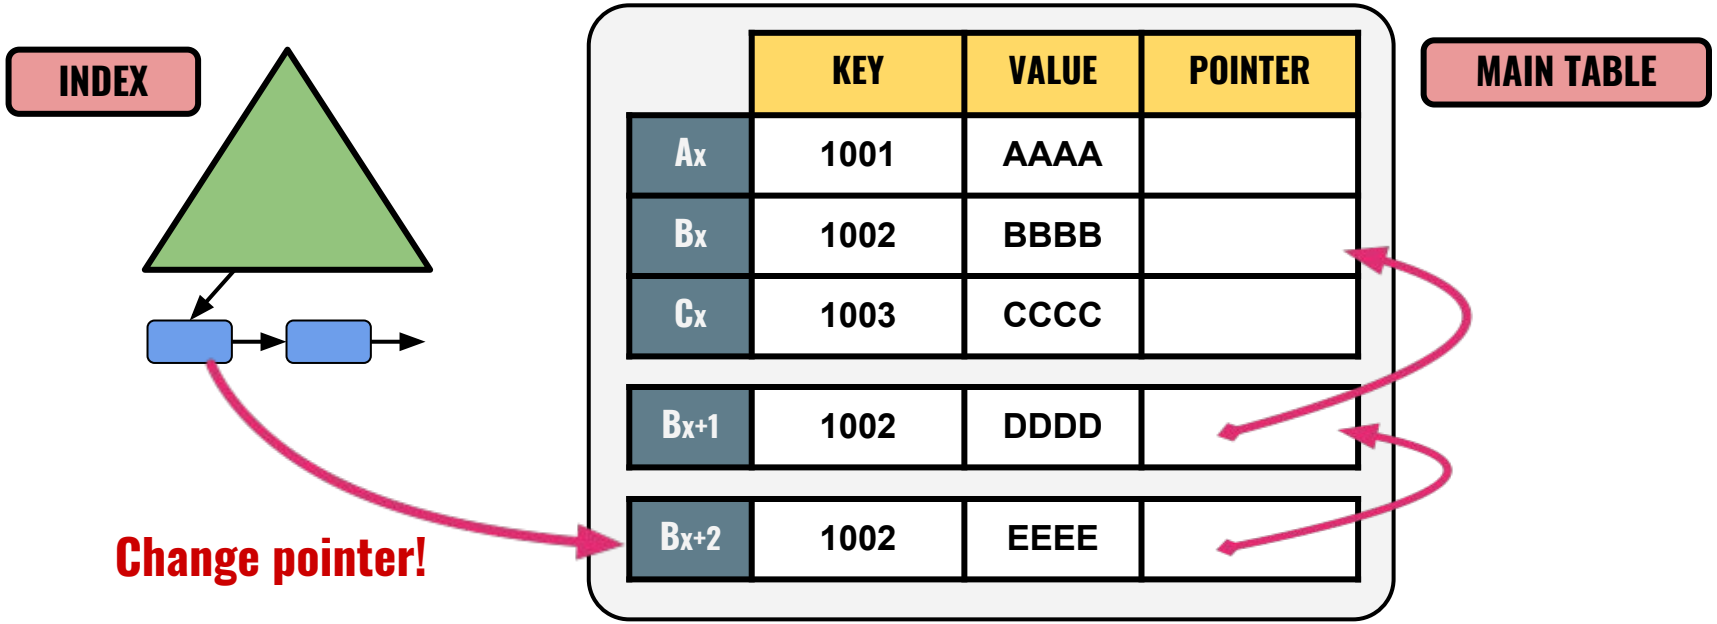

VERSION STORAGE

- **Approach #2: Time-Travel Storage**
  - Old versions are copied to separate table space.

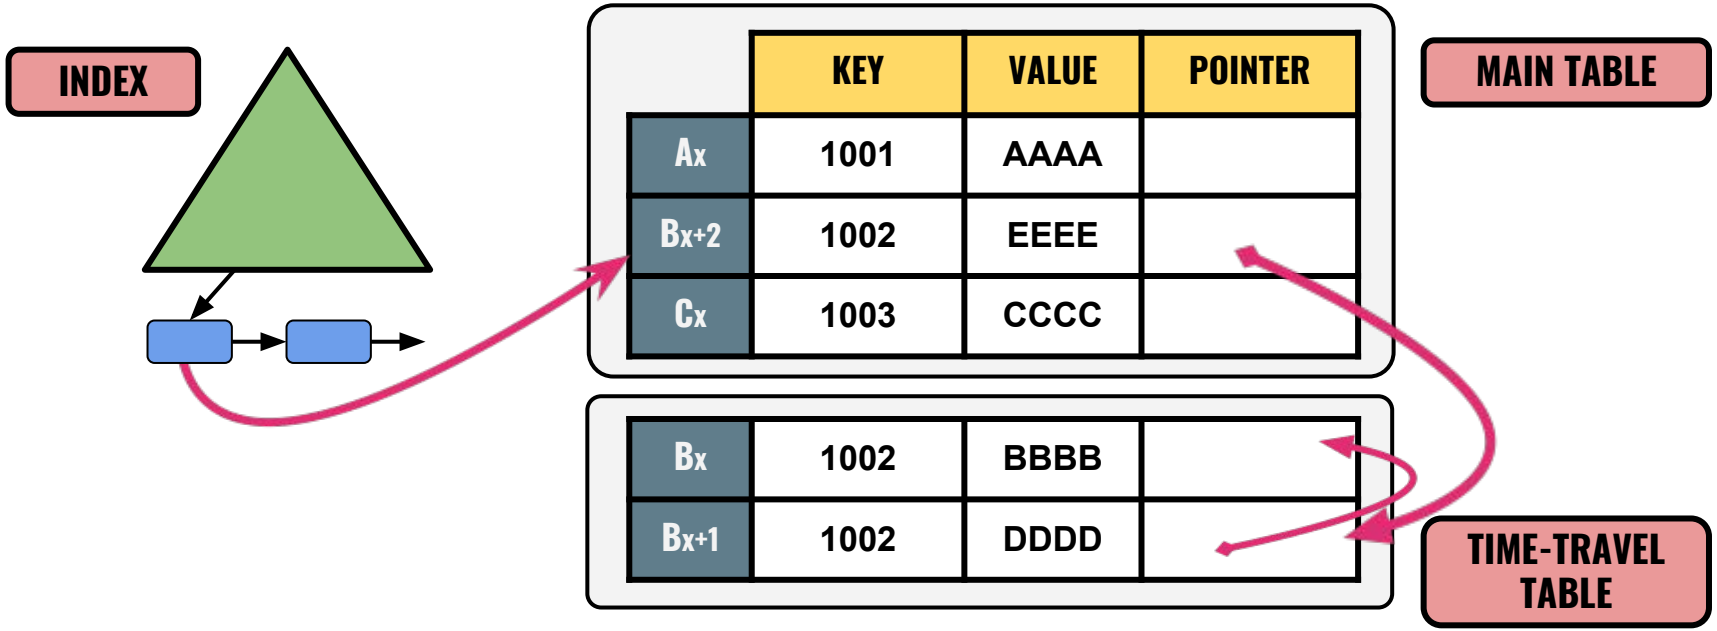

VERSION STORAGE

- **Approach #3: Delta Storage**
  - The original values of the modified attributes are copied into a separate delta space.

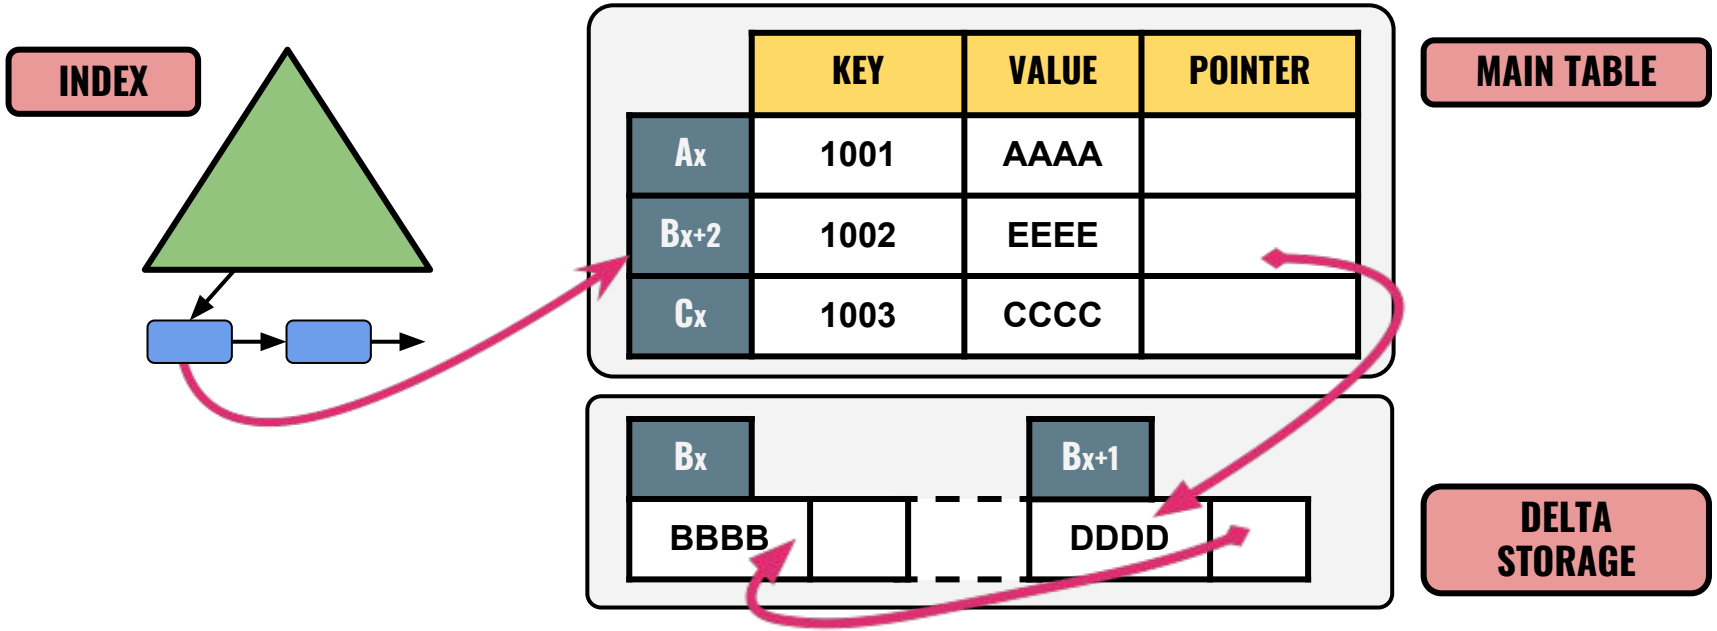

GARBAGE COLLECTION

SCHEME	DBMS
Tuple-Level	
Transaction-Level	

GARBAGE COLLECTION

- **Approach #1: Tuple-level**
  - Find old versions by examining tuples directly.

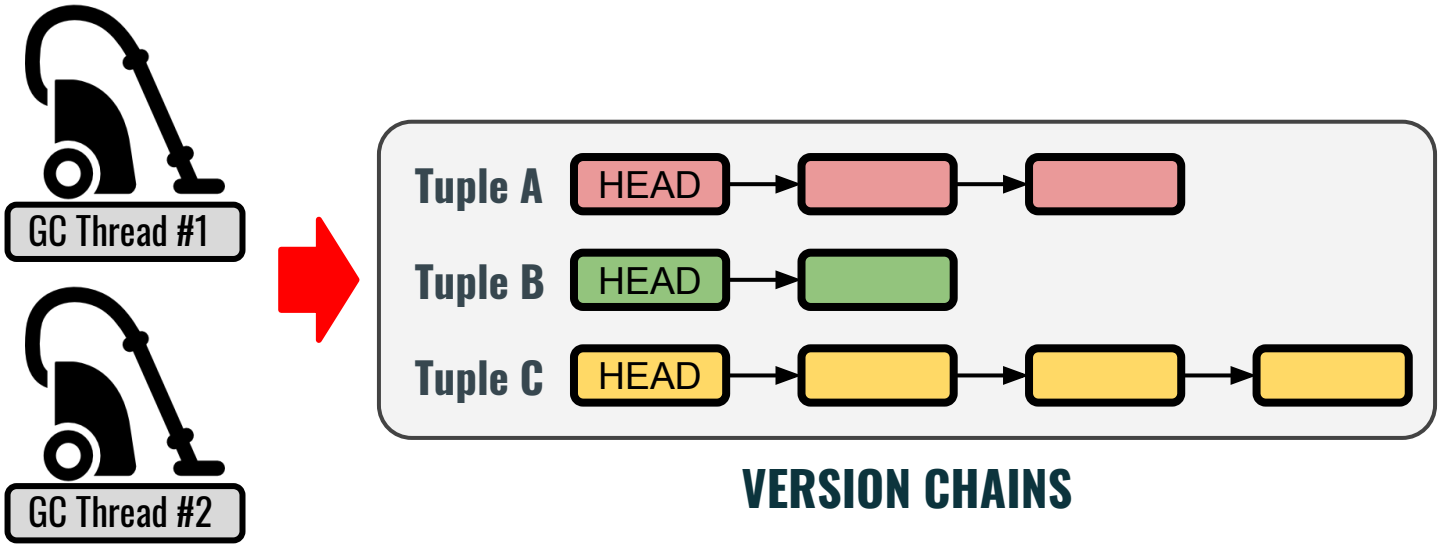

GARBAGE COLLECTION

- **Approach #2: Transaction-level**
  - Transactions keep track of their old versions so the DBMS does not have to scan tuples to determine visibility.

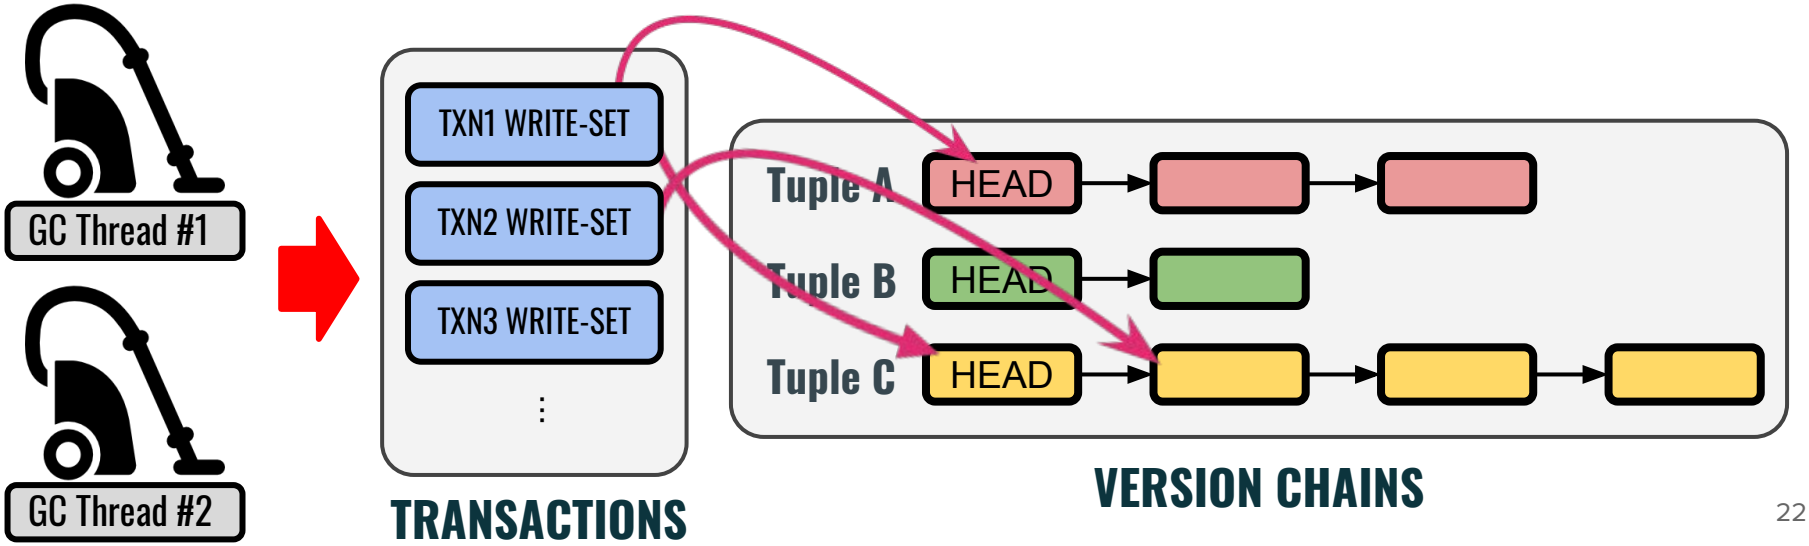

INDEX MANAGEMENT

SCHEME	DBMS
Logical Pointers	
Physical Pointers	

INDEX MANAGEMENT

- **Approach #1: Logical Pointers**
  - Use a fixed identifier per tuple that does not change.

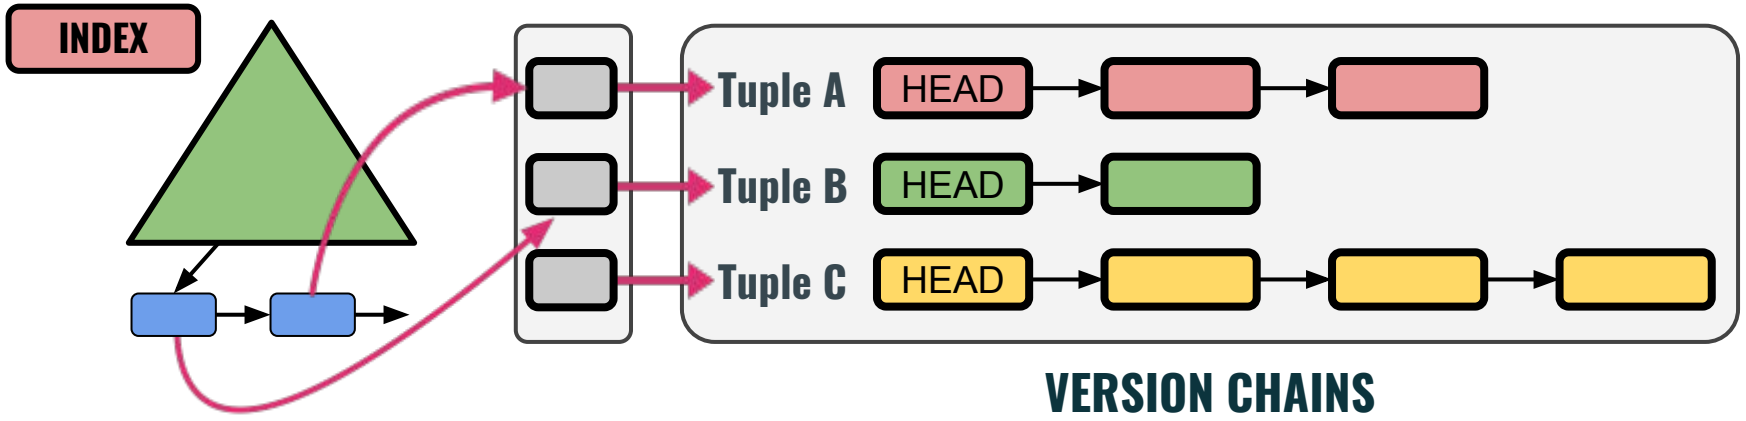

INDEX MANAGEMENT

- **Approach #2: Physical Pointers**
  - Use the physical address to the version chain head.

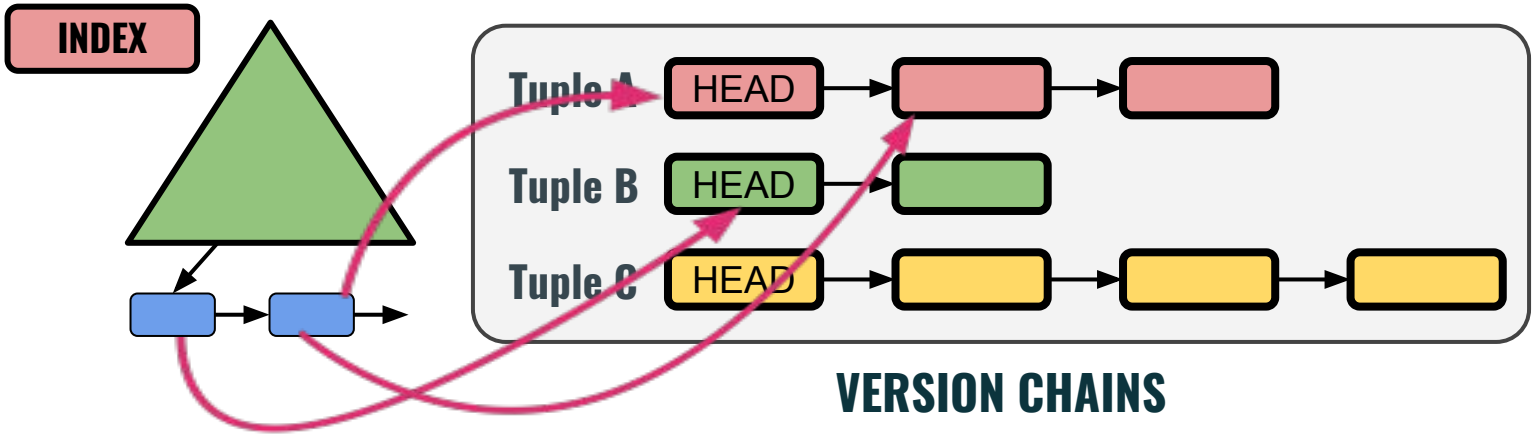

EVALUATION

- Benchmarks
  - YCSB
  - TPC-C
- Configuration
  - 4X Intel Xeon E7-4820 (40 cores)
  - 128 GB DRAM

EVALUATION

- Concurrency Control Protocol

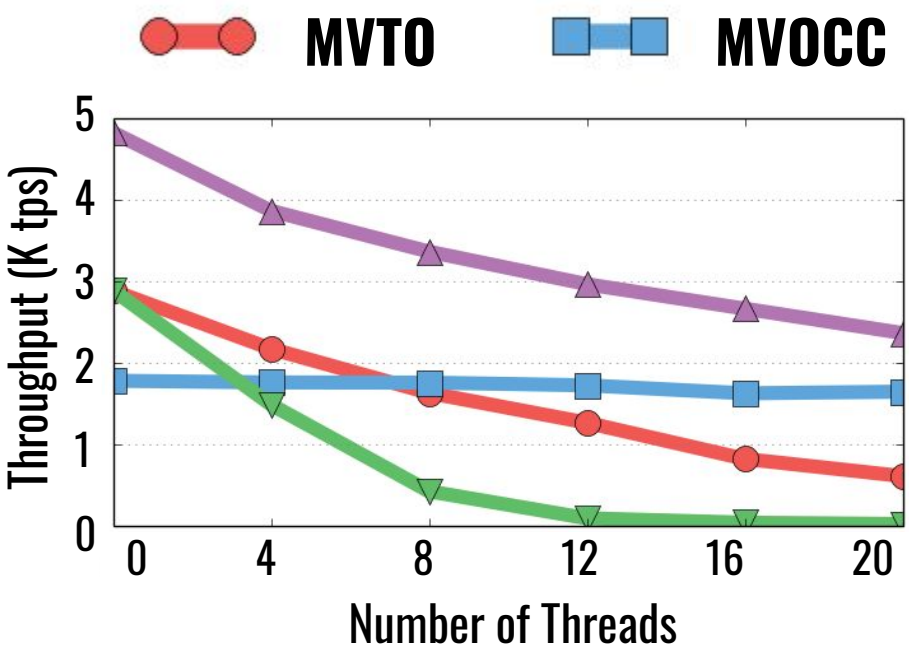

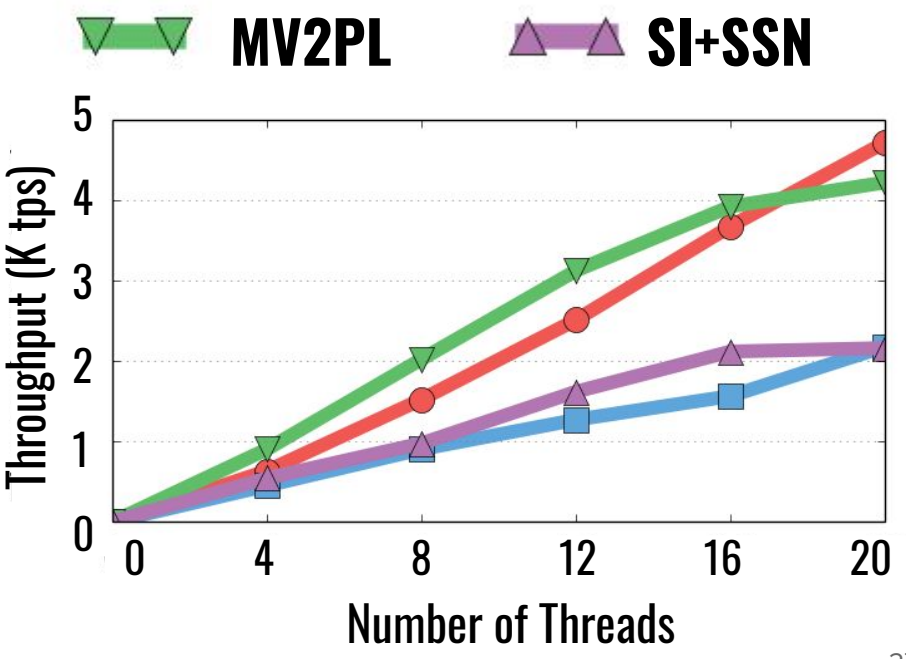

EVALUATION

- Version Storage

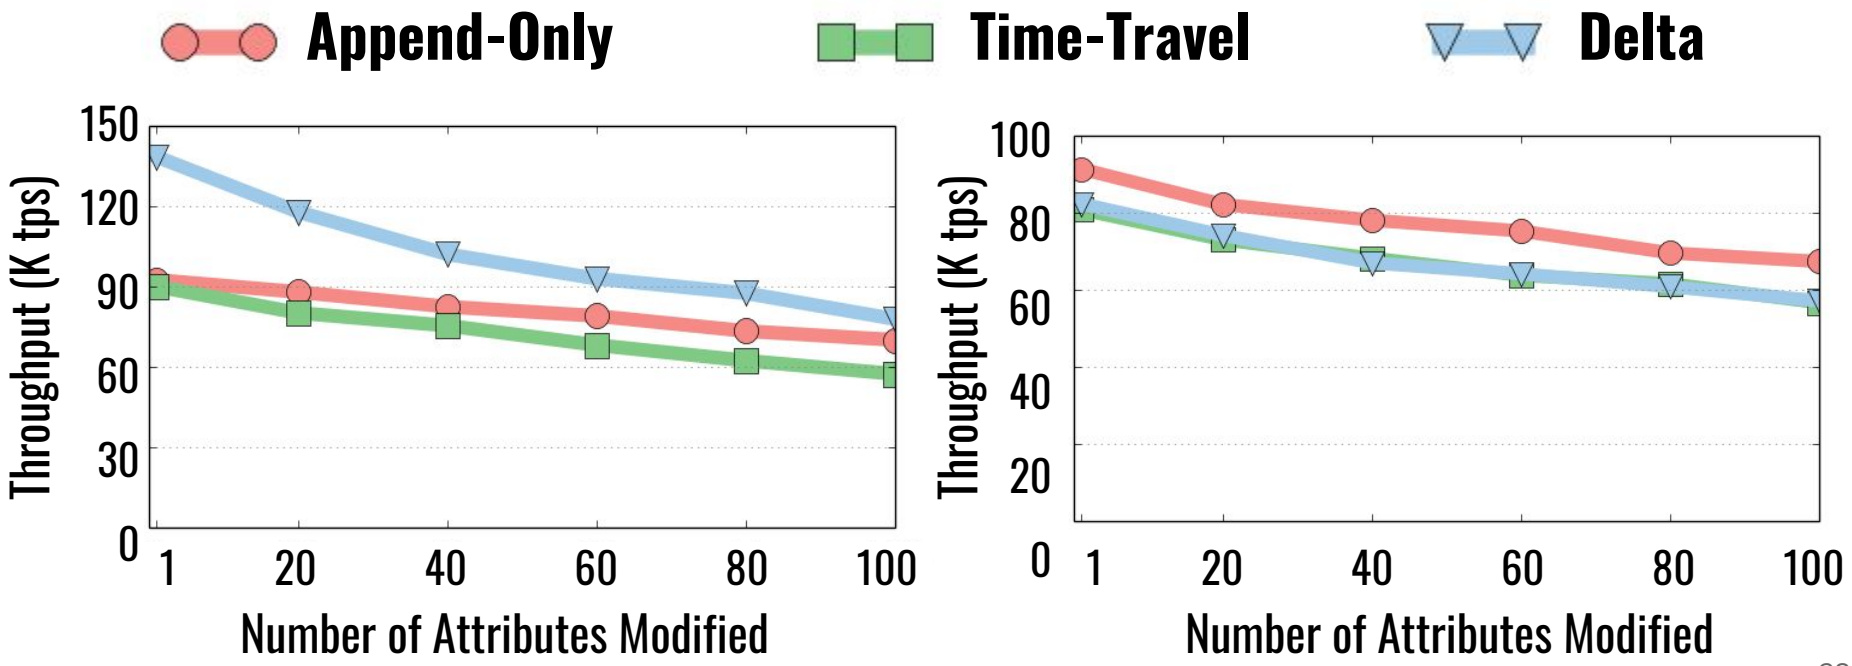

EVALUATION

- Garbage Collection

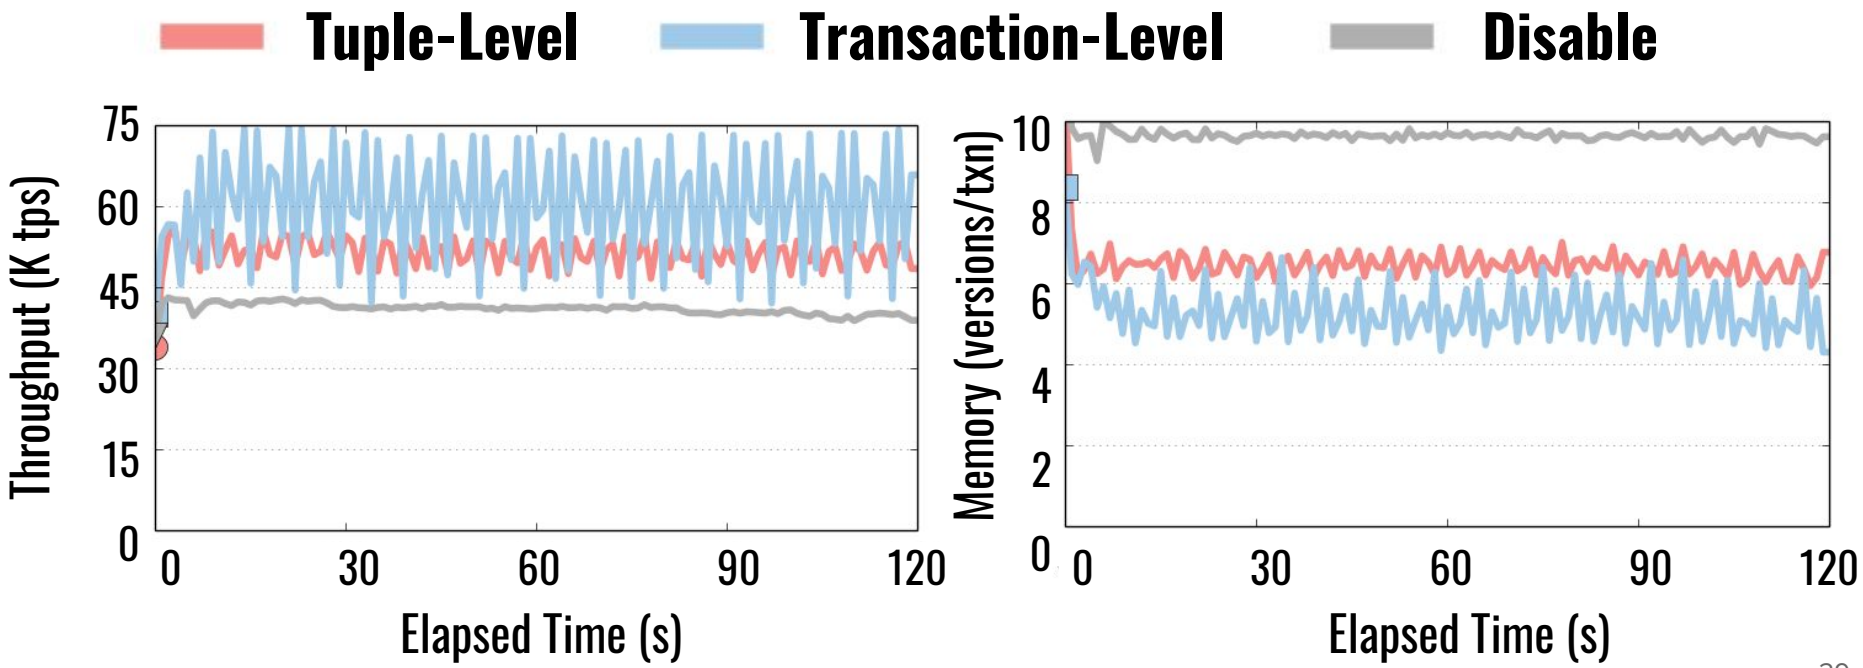

EVALUATION

- Index Management

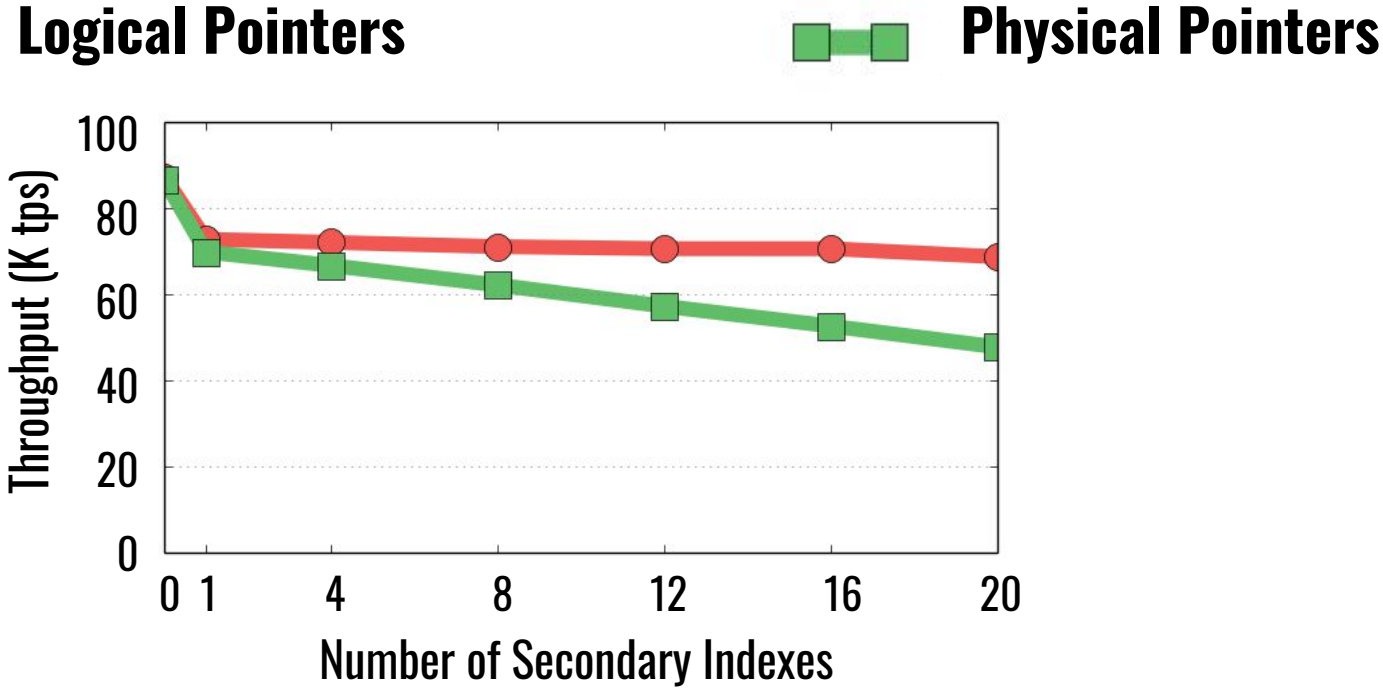

EVALUATION

- MVCC Configurations

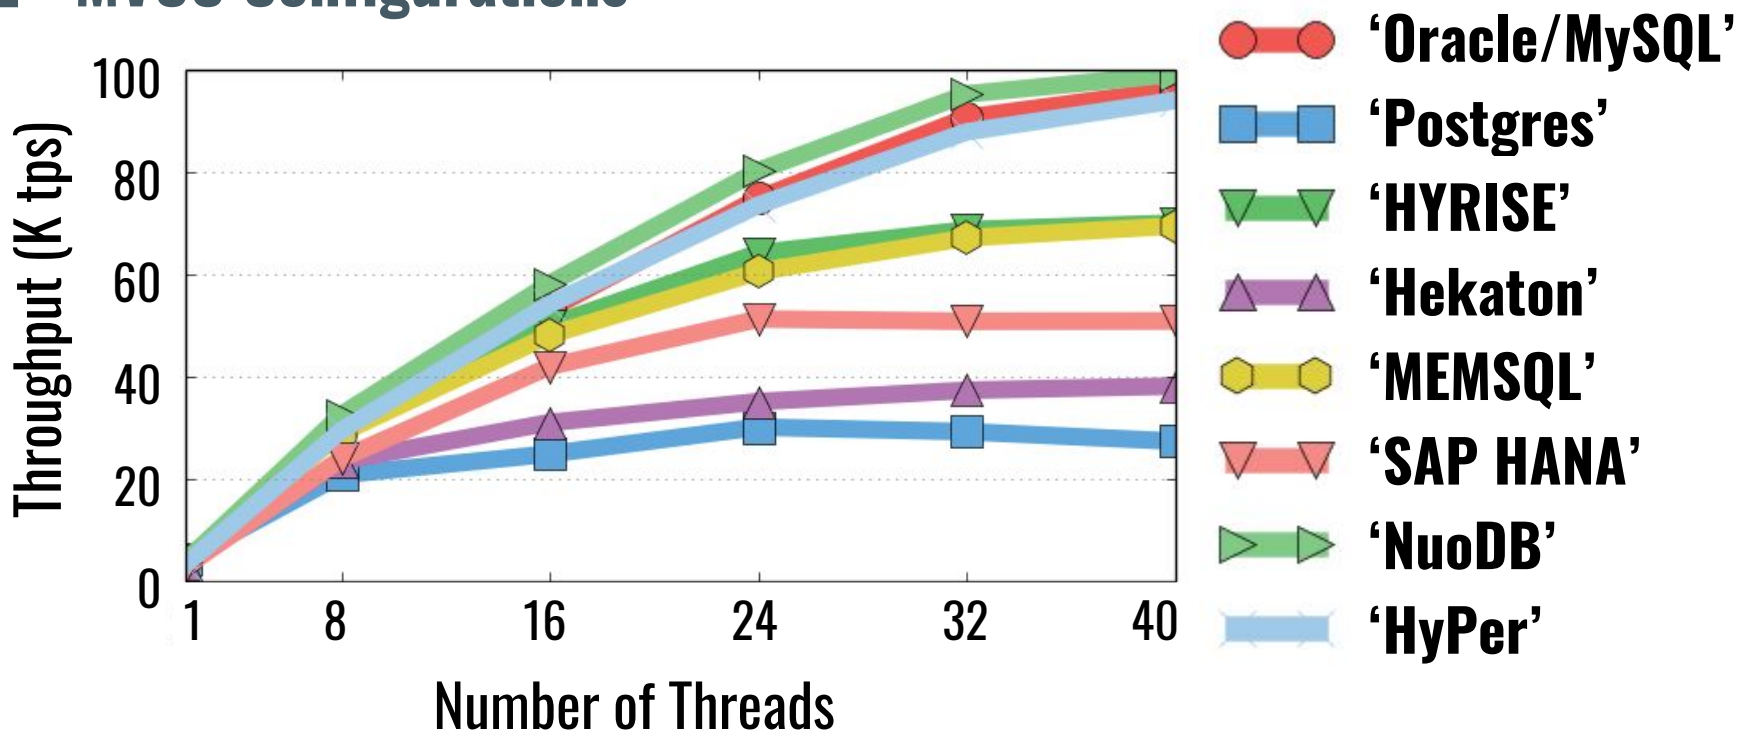

CONCLUSION

- Choosing the best MVCC scheme is challenging
  - Four design aspects
  - Multiple design decision combinations
  - Optimize for different objectives

END

@yingjunwu