

# Big Data Systems on Future Hardware

Bingsheng He

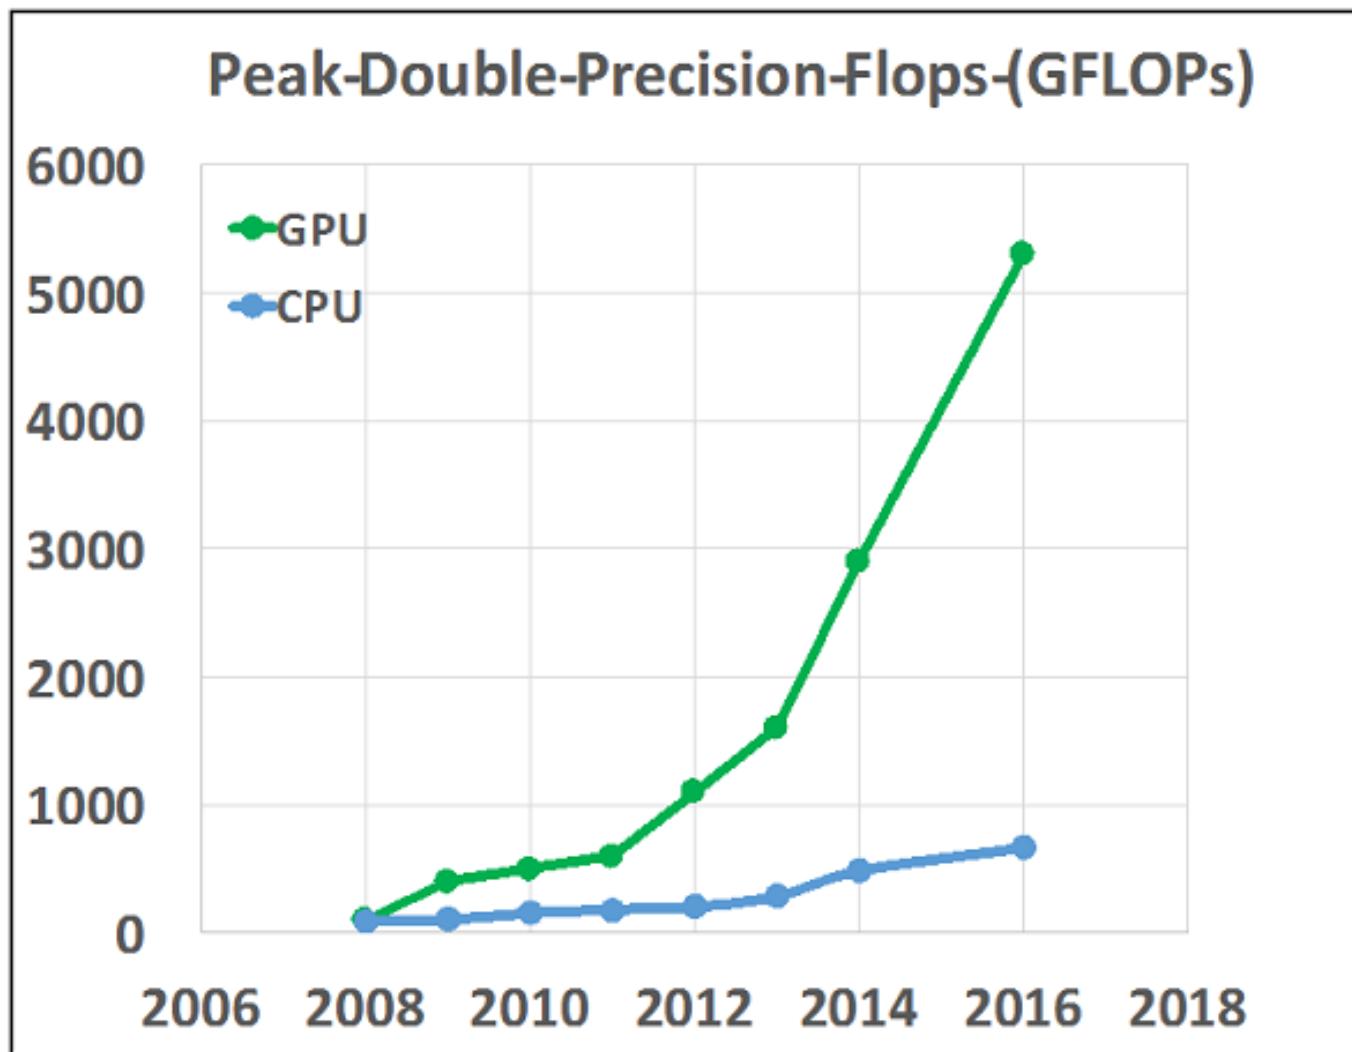
NUS Computing

<http://www.comp.nus.edu.sg/~hebs/>

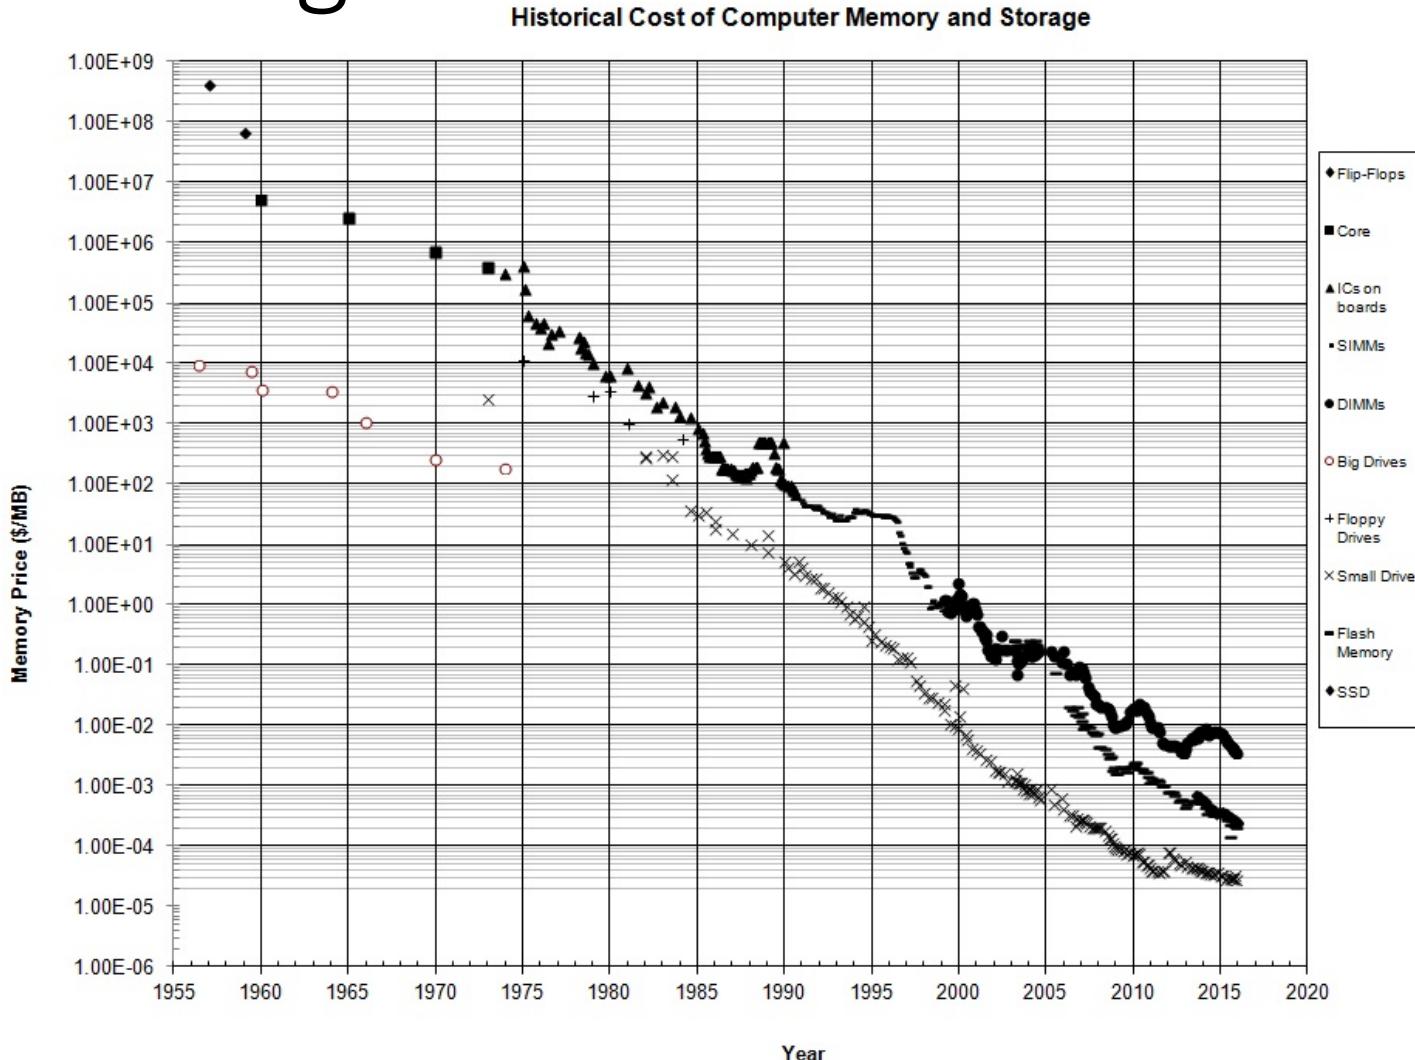
# Outline

- Why Hardware Matters?
- Challenges & Experiences
- On-going Projects
- Summary

# Why HW Matters: Processors

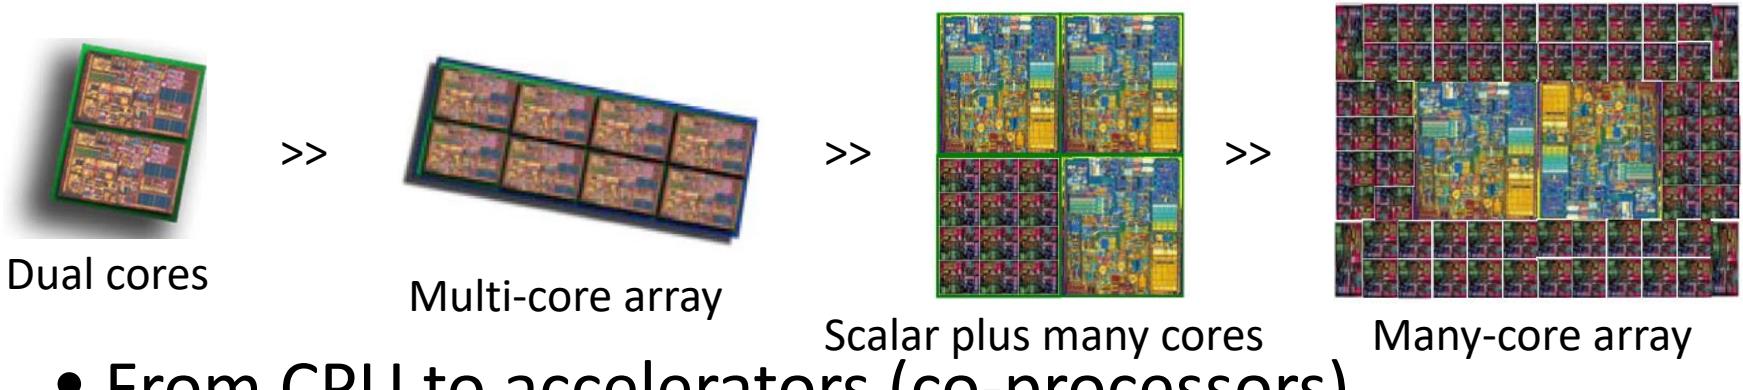


# Why HW Matters: Memory and Storage



# Emerging HPC Hardware: Parallelism and Heterogeneity

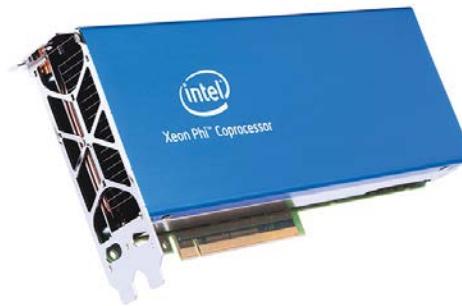
- Towards many cores



- From CPU to accelerators (co-processors)



GPU



Xeon Phi



FPGA

# Emerging HPC Hardware: Parallelism and Heterogeneity (Cont')

- Towards tightly coupled heterogeneous systems



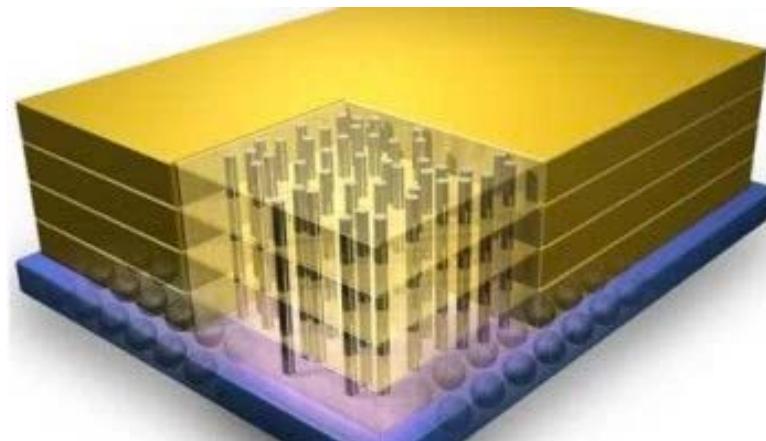
AMD APU



Intel-Altera Heterogeneous Accelerators

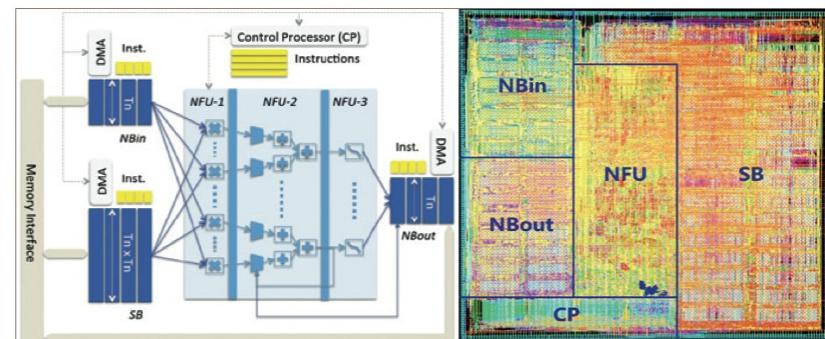
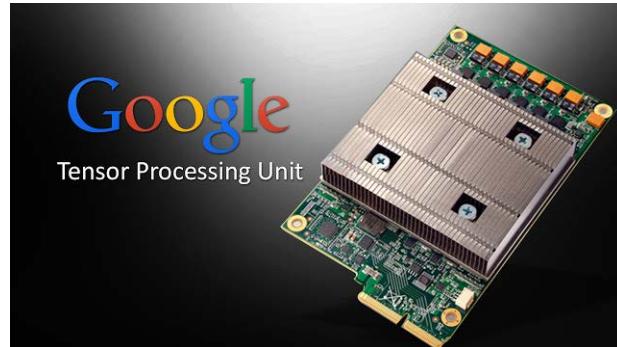
...

- High bandwidth memory (HBM)



# Emerging HPC Hardware: Generalization vs. Specialization

- AI chips
- Graph accelerators
- ...

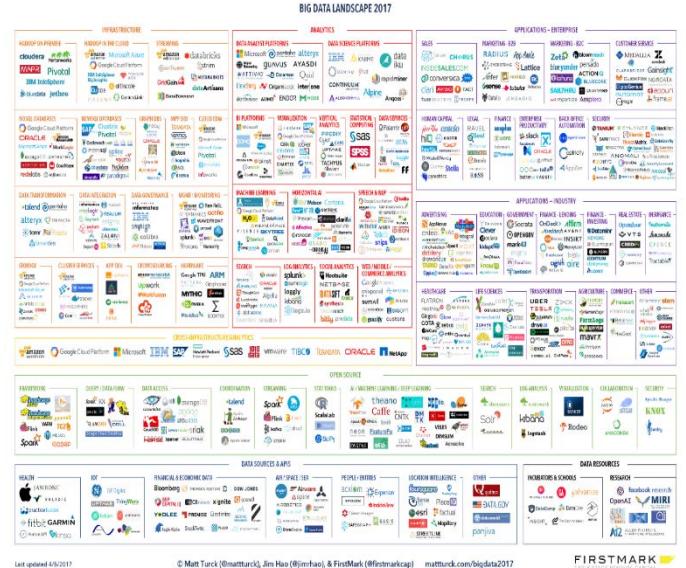
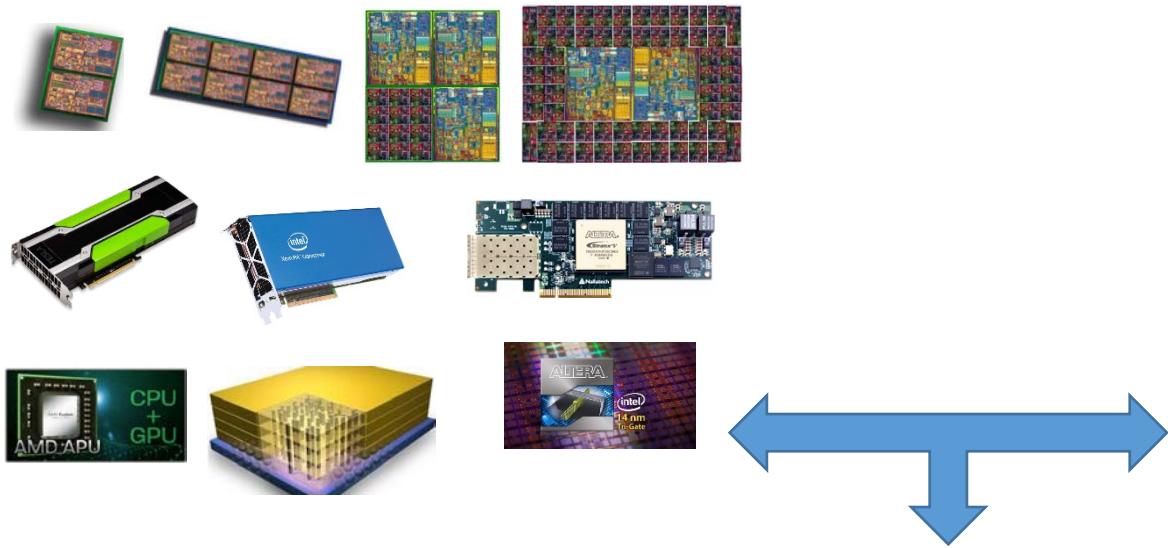


# Future Hardware

- Processors
  - 1, 000 cores
  - Heterogeneous/specialized hardware (FPGA/ASIC)
- Disk is dead, NVRAM is disk, DRAM is cache,  
Locality is still the King.
  - NVRAM/3D stacking
  - “Tape is Dead, Disk is Tape, Flash is Disk, RAM Locality is King” by Jim Gray
- Cluster as a personal supercomputer
  - Fast and cheaper interconnects  
(e.g., Infiniband)



# When Big Data Meets Emerging Hardware (Con't)



- **Architecture-aware optimizations:** how to make our software systems fully optimize for the hardware?
- **New hardware:** how can new applications drive the design of new hardware?
- **New systems/applications:** can superb hardware power enable new systems/applications?

# Outline

- Why Hardware Matters?
- Challenges & Experiences
- On-going Projects
- Summary

# Challenges & Experiences

- Challenge #1: HW is simply one side of the coin; SW is the other side.
- Challenge #2: The leap from prototype to production.
- Challenge #3: Millions of lines of legacy code.
- Challenge #4: Generalization vs. Specialization

# Challenge #1: HW is simply one side of the coin; SW is the other side.

- Parallel algorithmic design
  - Revisit existing parallel algorithm and develop new algorithm.
  - Running a large task in parallel vs. running many small tasks in parallel
- Hardware features
  - GPU: shared memory, cache, constant memory etc.
  - FPGA: local memory, pipeline, pipe etc.
- Architecture-aware software system design can make a big difference.

# Our Experiences in GPGPU-based Data Management Systems

CUDA was released in Feb. 2007	GPUQP (GDB) published in SIGMOD 2008 (“best papers”)	Mars (GPU-based MapReduce) published in PACT 2008 (2 <sup>nd</sup> top cited paper in PACT)*	Mars has been extended to AMD GPU and Hadoop (TPDS10)
OmniDB: relational database on coupled CPU/GPU architectures (VLDB’13/14/15, SIGMOD 16, VLDB’13 demo,...)	Medusa: GPU-based graph processing (TPDS13/14, VLDB13 best demo, CloudCom13)	Transaction executions on GDB (VLDB11)	GDB supports compressed column-based processing (VLDB10)
<p style="color: red;"><b>Example speedup over CPU:</b></p> <ul style="list-style-type: none"><li>• OLAP: 2-7X (vs. one CPU)</li><li>• OLTP: 4-10X (vs. one CPU)</li><li>• Graph: 5.5X (vs. two CPUs)</li></ul>			14

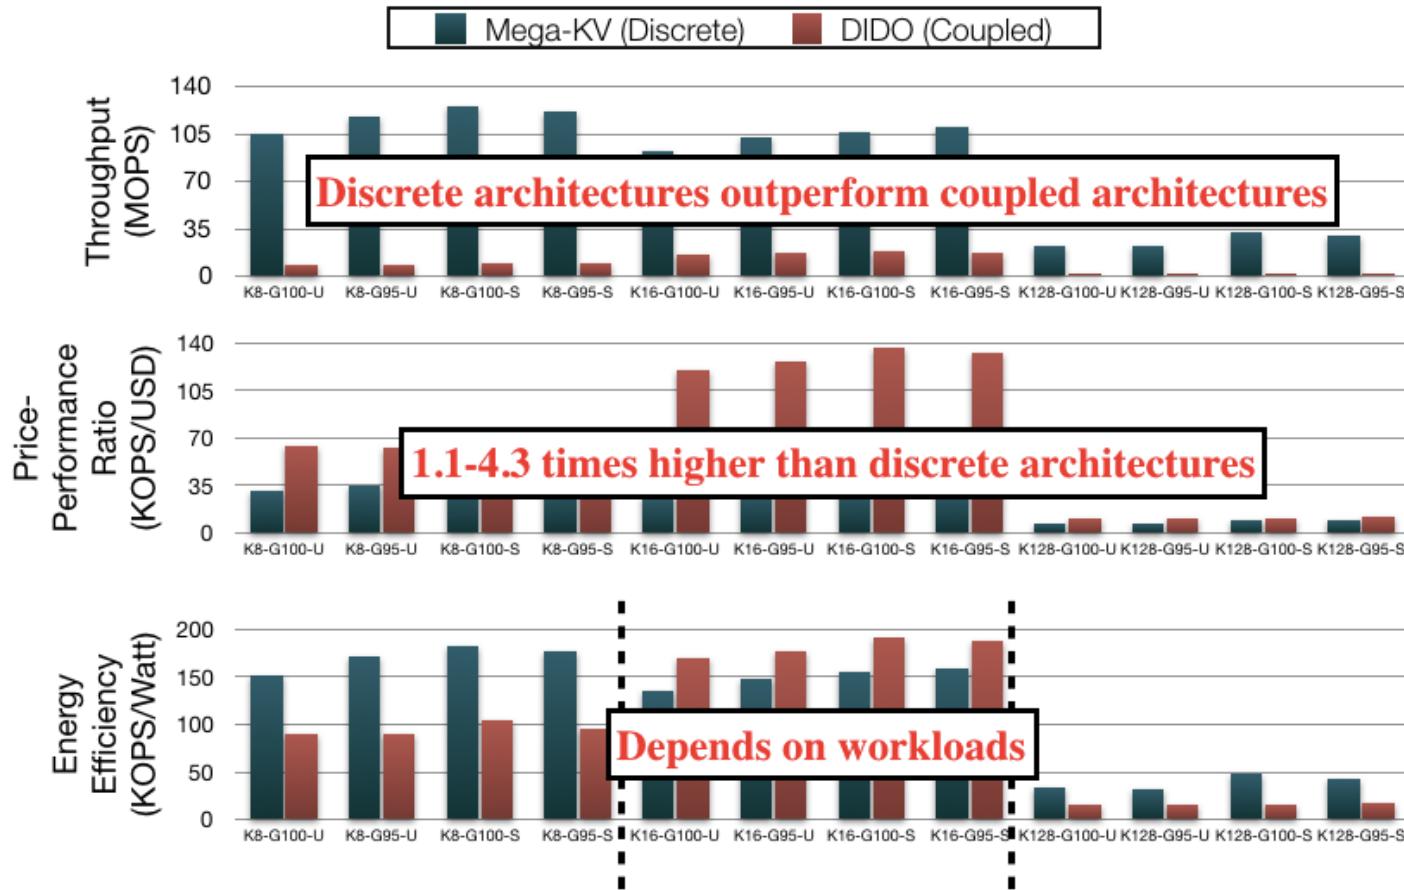
# Challenge #2: The leap from prototype to production

- Many research papers demonstrate that the speedup of GPU/FPGA can reach 10-100X.
- Still, GPU/FPGA has rather limited adoptions in production environments (although increasingly more).
- Why?
  - Hardware: power supply budget, space, costs,...
  - Software: software maintenance, reliability, manpower expertise, sharing, virtualization ....
  - Workload: deep learning, database ...
- Besides algorithmic innovation, various system aspects have to be addressed.

# Experiences in Addressing Practical Issues of GPU Computing

- PCI-e bus via data compressions
  - Database compressions on column-based databases [VLDB2010]
- Concurrent kernel executions
  - Resource complementary kernel co-scheduling [TPDS 2014]
- GPU virtualizations
  - Gaming virtualization [USENIX ATC 2016]
- Minimizing data flow overhead among processors
  - Pipeline execution [SIGMOD 2016]

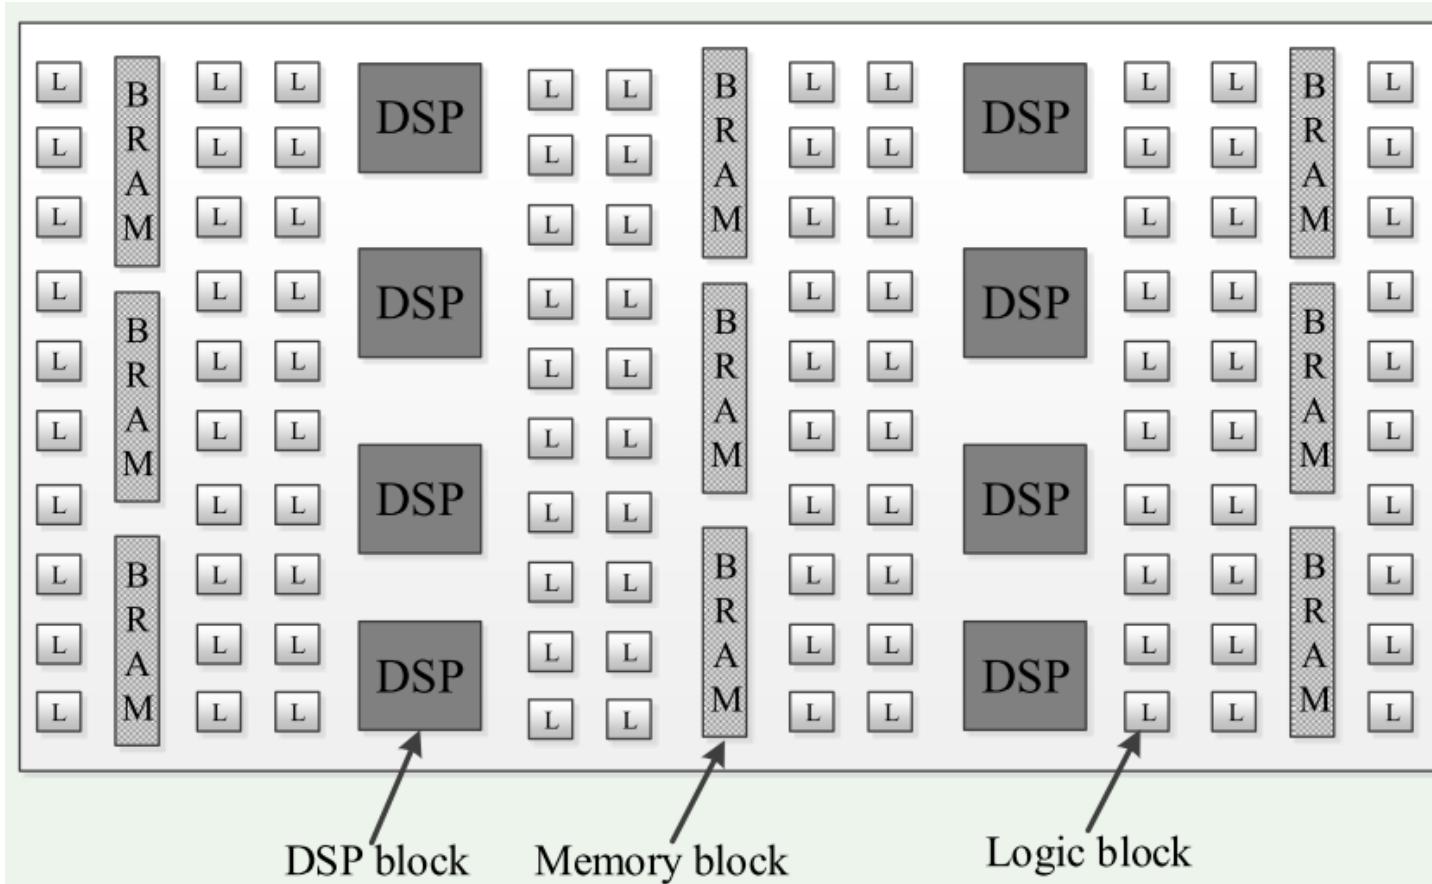
# Performance, or Perf per \$, or Perf per Joule?



# Challenge #3: Millions of lines of legacy code

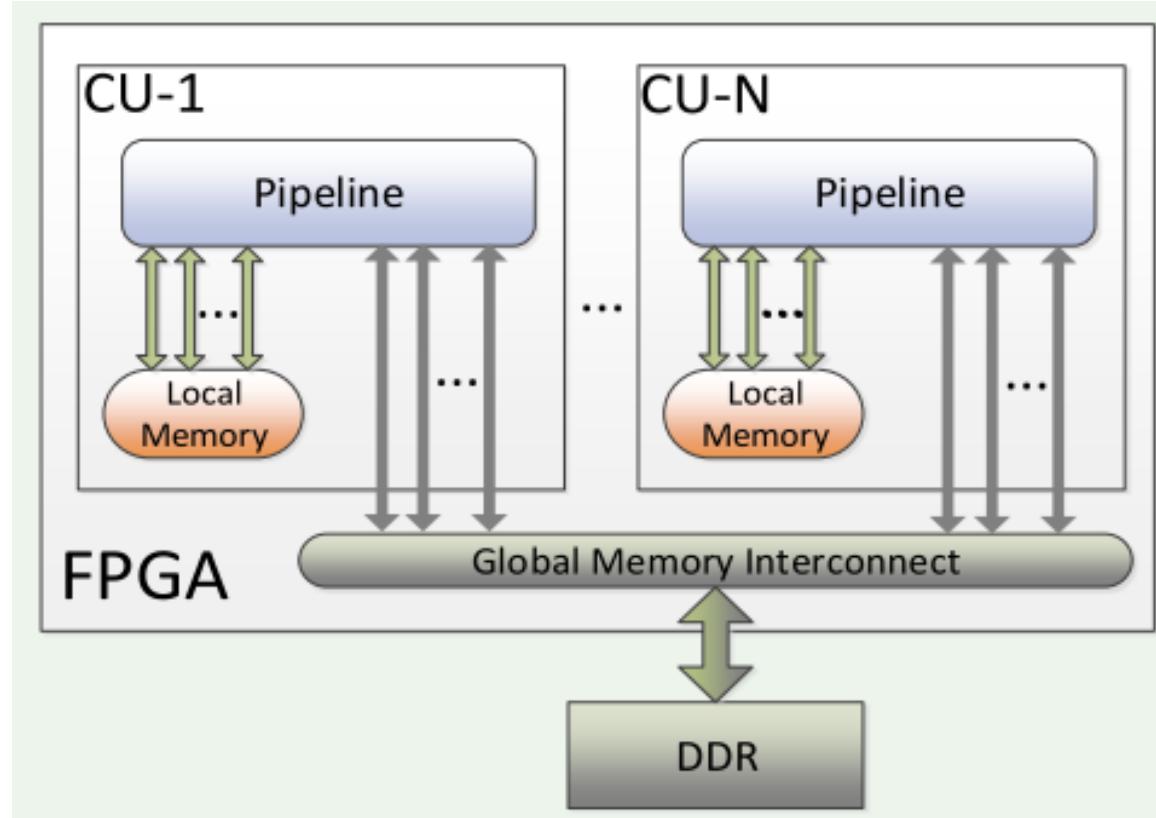
- Our legacy software systems are monsters
  - National labs have MPI programs of millions of code lines.
  - Google's Internet services spans some 2 billion lines of code.
  - Microsoft's Windows operating system has around 50 million lines.
  - Other younger ones: Hadoop 2millions, Spark 0.9 million, MySQL 2.7 millions...
- The reality is, “write once, reuse till many many times”.
- The research on automatic parallel optimizer is dead.
- (Semi-)Automated tools are needed to resolve the pain points.

# “Architectural Evolution” of FPGA (Field Programmable Gate Arrays)



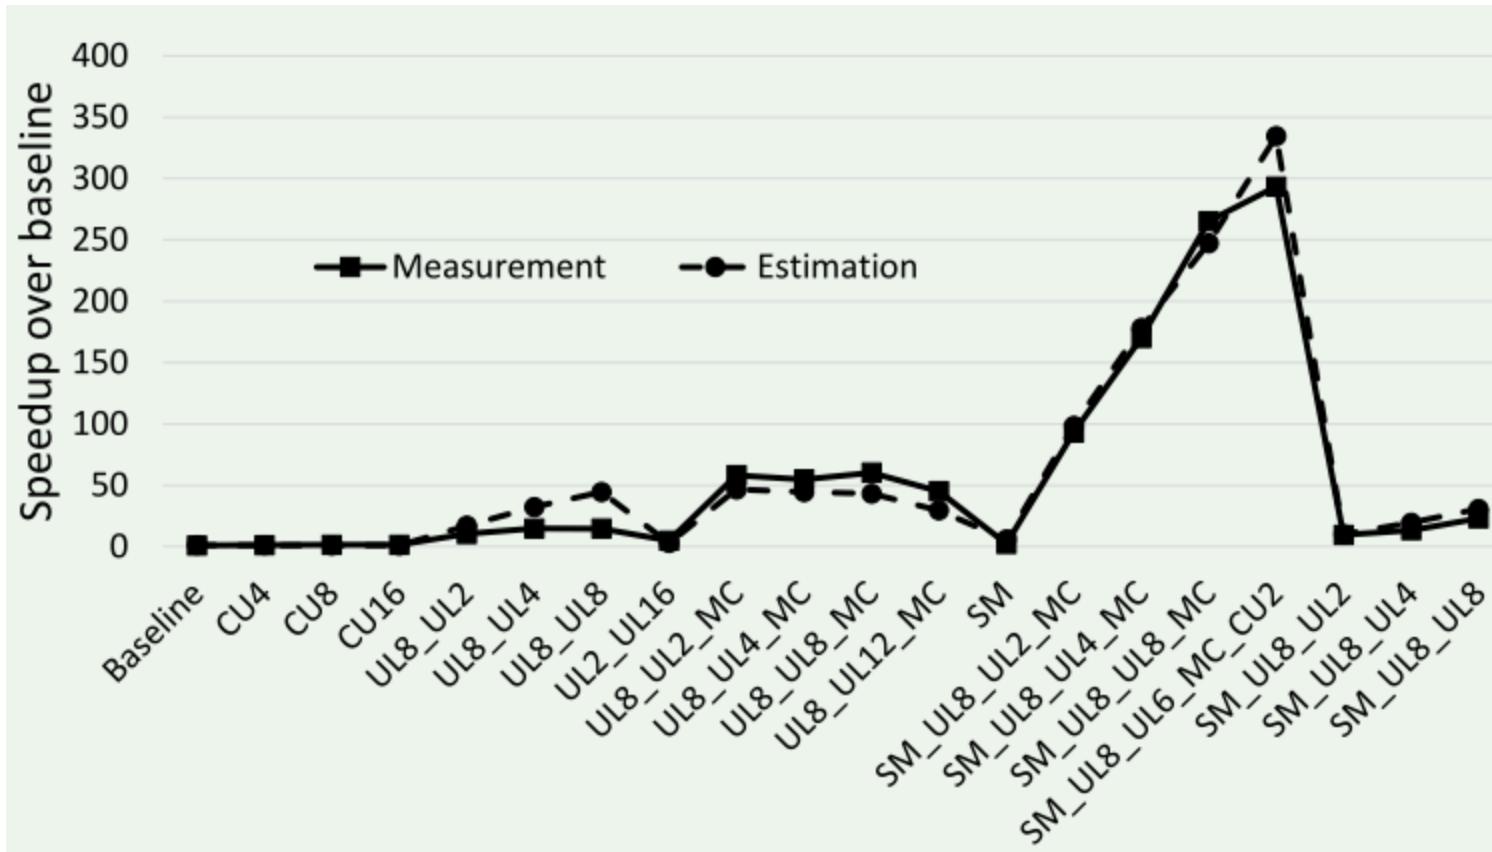
- **Hardware centric**
- **Users need to program with low-level hardware description languages.** ☹

# “Architectural Evolution” of FPGAs: From OpenCL’s Perspective



- Software centric → FPGA is viewed as a parallel architecture.
- Users can program with OpenCL. ☺

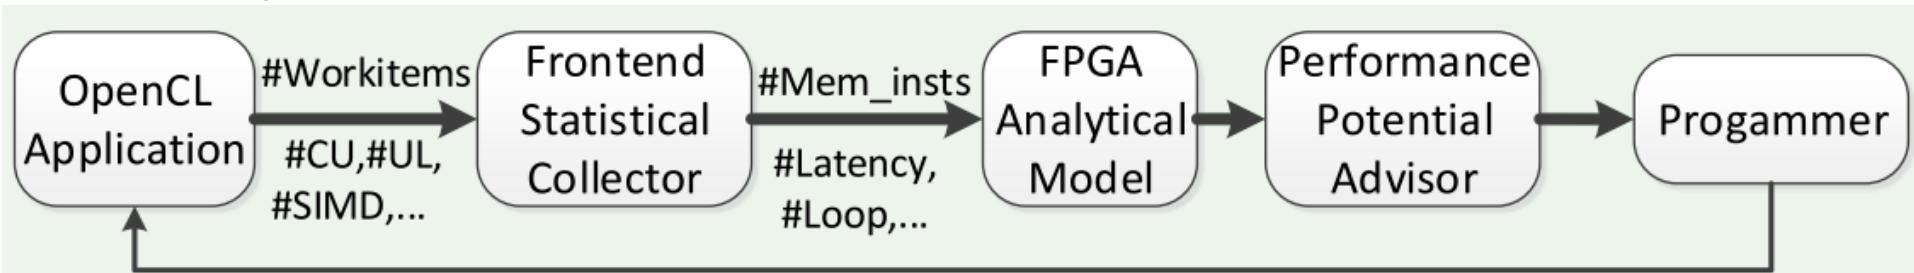
# Our Experiences on FPGA



- The example: K-Means
- Applying different combinations of optimization leads to huge performance differences → Tools for optimizations are needed.

# Our Solution: Static and Dynamic Program Analysis

- We propose a performance analysis framework to assist programmers to optimize the OpenCL program on FPGA
  - Static statistical collection on the corresponding LLVM IR code.
  - Dynamic profiling of the OpenCL application execution.
  - FPGA analytical model predicts the performance of OpenCL application.
  - The performance advisor digests the model information and provides the four potential metrics to understand the performance bottleneck.



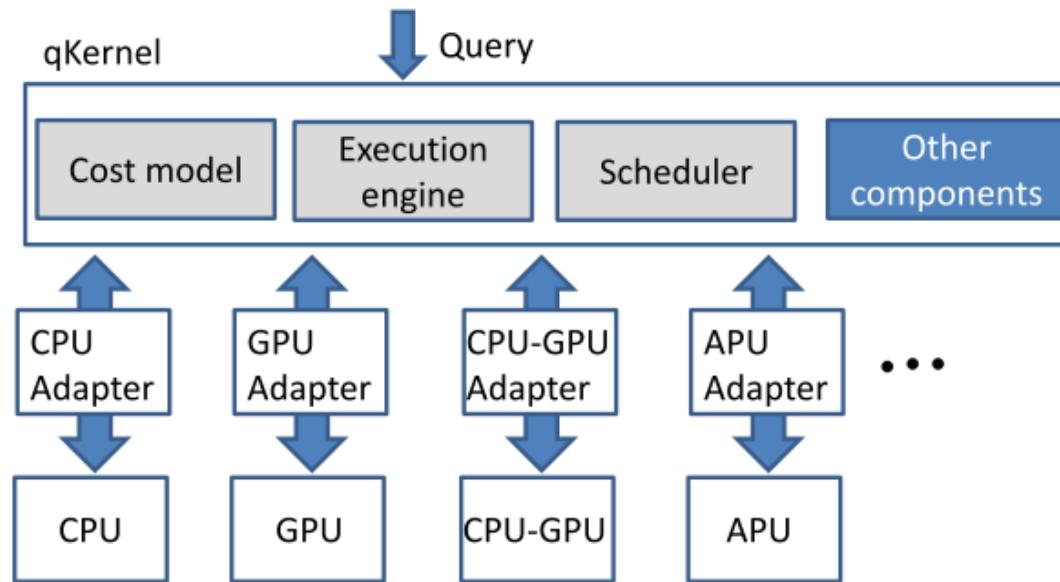
Details in “Zeke Wang<sup>A</sup>, Bingsheng He, Wei Zhang, Shunning Jiang. A Performance Analysis Framework for Optimizing OpenCL Applications on FPGAs. HPCA 2016”

# Challenge #4: Generalization vs. Specialization

- Specialized hardware
  - “SQL in Silicon” (in Oracle SPARC M7 processor)
  - Google TPU for deep learning
- Specialized software
  - System call overhead for memcached
  - Layers of abstractions in OS (e.g., for NVRAM)
- A compromise is possible (but difficult to find the optimal cut between generalization and specialization).

# SW Portability vs. Specialization

- OmniDB: General Engine Design + Adapter to Specific Architecture



# Outline

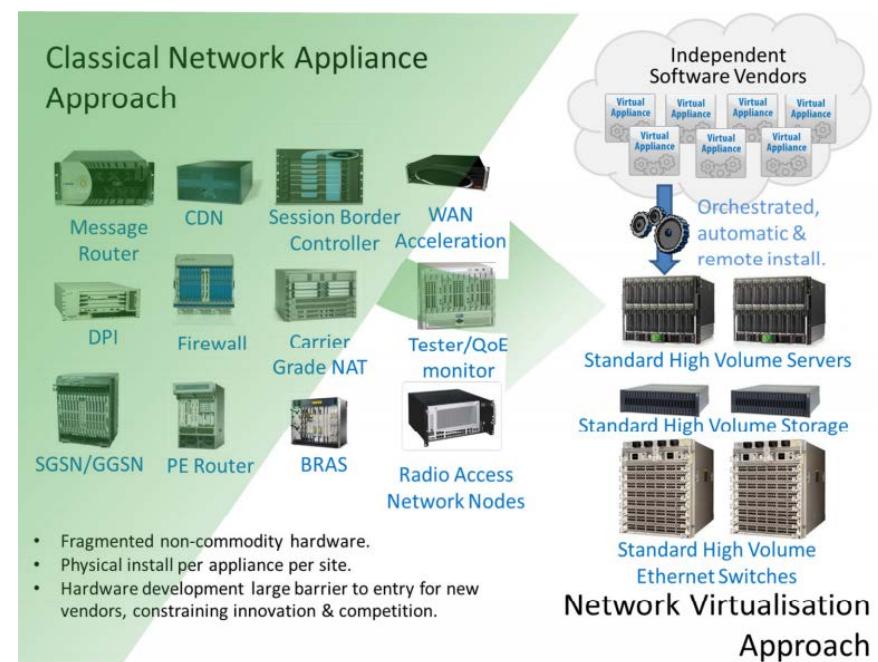
- Why Hardware Matters?
- Challenges & Experiences
- On-going Projects
- Summary

# On-going Projects

- High performance graph databases
  - Goal: a graph database that can outperform Neo4j by 10+times.
- Network Function Virtualization (NFV)
  - Goal: a cost-effective solution to replace specialized network devices.

# Network Function Virtualization

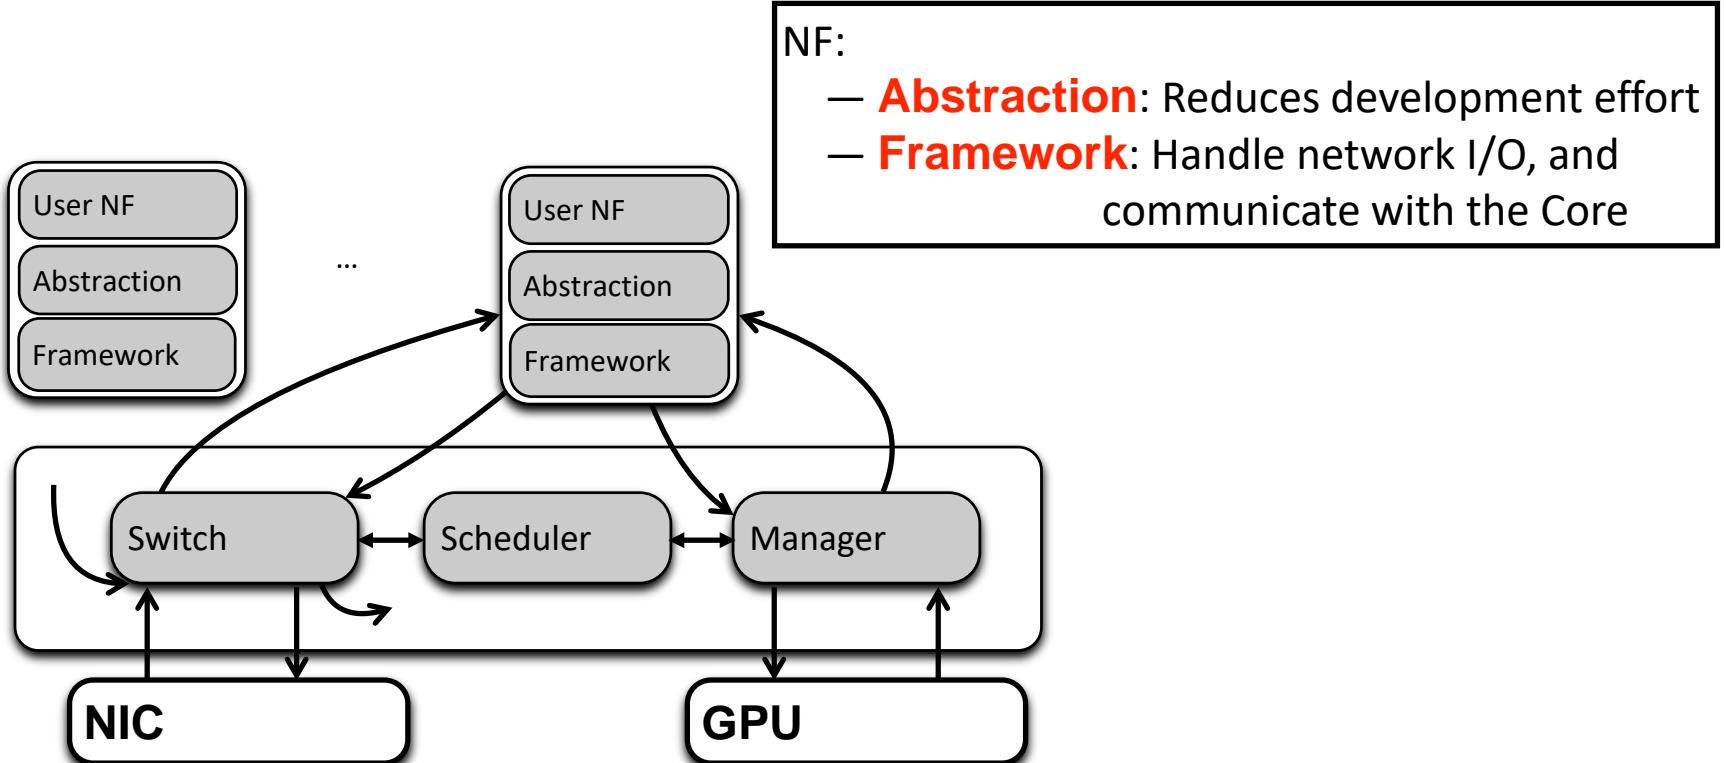
- Basic trend: supporting network functions shifting from specialized hardware to software.
- Huawei predicts, in 3-5 years, CPU + GPU will satisfy the computational requirement for telecom equipment.
- Challenges:
  - Massive network flows
  - Low latency
  - Flexible deployment



# GPU-Accelerated NFVs?

- GPU advantage over specialized hardware
  - Comparable speed
  - Low cost
- However,
  - Low GPU Utilization (Each VM has no enough traffic to fully utilize GPUs) → GPU virtualization is still an open research problem.
  - Unpredictable Latency
  - Lots of development effort

# Poseidon: Use Remote API to Virtualize GPUs in NFV



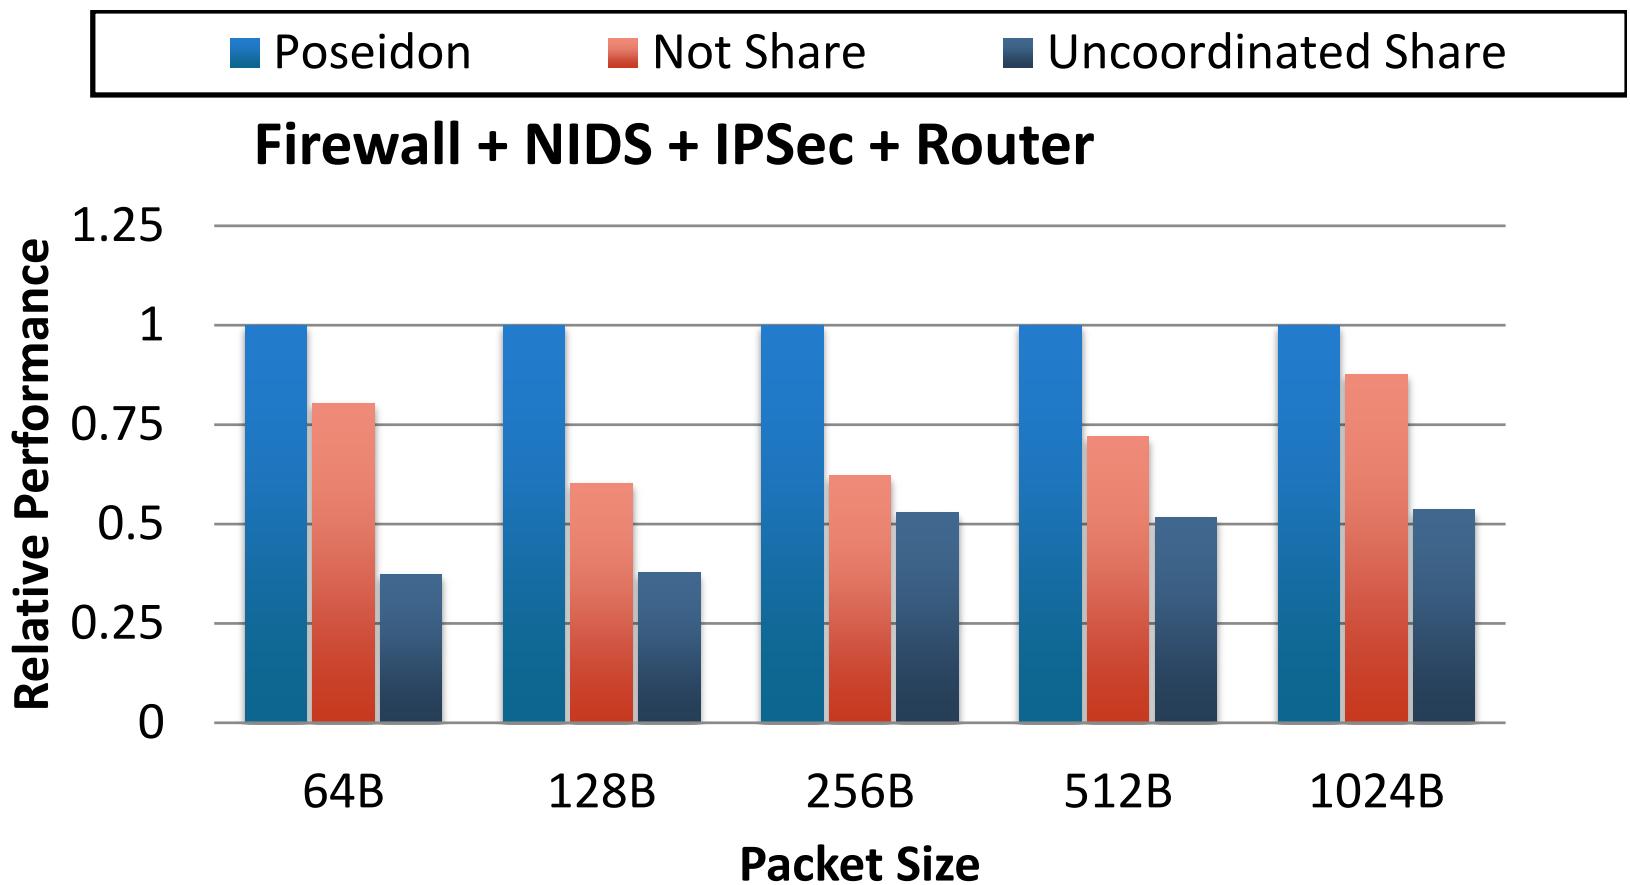
Core:

- **Switch**: Handle network I/O
- **Scheduler**: Schedule GPU for execution
- **Manager**: Launch GPU kernels

NF:

- **Abstraction**: Reduces development effort
- **Framework**: Handle network I/O, and communicate with the Core

# Preliminary Results



Our system can significantly improve the performance of GPU sharing.

# Outline

- Why Hardware Matters?
- Challenges & Experiences
- On-going Projects
- Summary

# Summary

- (Database) systems on emerging hardware continue to be a challenging and exciting research area.
- Our experiences demonstrate the system insights as well as open challenges of building big data systems on future architectures.
- Hardware and software co-design might be the key for the success of this battle.

# Thank you!

More about Xtra Computing Group:

<http://www.comp.nus.edu.sg/~hebs/>