

Shuhao Zhang

Professor, School of Computer Science and Technology, HUST

shuhao_zhang@hust.edu.cn ORCID: 0000-0002-9927-6925 home.shuhao.sage.org.ai

Profile

Shuhao Zhang is a Professor in the School of Computer Science and Technology at Huazhong University of Science and Technology. He leads a task under a major national science and technology project of the Ministry of Science and Technology of China and is a recipient of the National Natural Science Foundation of China Excellent Young Scientists Fund (Overseas). In 2026, he was awarded support through the NSFC Research Fund for International Excellent Young Scientists (RFIS-II).

His research aims to build high-performance inference engines for domestic accelerators in China, forming the core of continuous data and inference systems for agent applications. With vLLM-HUST as an engineering vehicle, he studies scheduling, KV-state management, execution and communication optimization. Orchestration, external memory, dynamic data maintenance, and evaluation components extend this core toward long-running agents.

Appointments

Huazhong University of Science and Technology

Professor

School of Computer Science and Technology.

Nanyang Technological University

Assistant Professor

Before joining HUST.

Technische Universität Berlin

Postdoctoral Researcher

Postdoctoral research on data systems and distributed stream processing.

Education

National University of Singapore

Ph.D., Computer Science, 2019

Research area: database systems, stream processing, parallel and distributed systems.

Nanyang Technological University

B.Eng., Computer Engineering, 2014

Research Focus

- Inference engine core: scheduling, KV-state reuse, execution and communication on domestic accelerators.
- Extensions: dataflow orchestration, external memory, dynamic retrieval and graph maintenance.
- Validation: reproducible workloads, end-to-end serving quality, and agent applications.

Major Projects and Funding

- Principal Investigator, NSFC Research Fund for International Excellent Young Scientists (RFIS-II), awarded in 2026.
- Task Lead, major national science and technology project of the Ministry of Science and Technology of China.
- Principal Investigator, National Natural Science Foundation of China Excellent Young Scientists Fund (Overseas).
- Principal Investigator and collaborator on previous research grants from Singapore MOE, Singapore NRF, Temasek Laboratories, and university-level research programs.

Representative Systems

- **vLLM-HUST**: vLLM-based research and engineering stack for inference on domestic accelerators; the core of the systems agenda.
- **SAGE**: Dataflow orchestration extending the inference core toward modular, controllable LLM reasoning.
- **Neuromem**: External-memory lifecycle research and benchmarks; part of the continuous-state extensions.
- **BriskSeed / FlowRAG / StreamFP / GRACE / CANDOR-Bench**: Online search-history reuse, dynamic retrieval, stream learning, graph maintenance, and continuous-ANNS evaluation.
- **MorphStream / TStream / CStream**: Stream processing systems for multicore, stateful, and hardware-conscious execution.

Selected Recent Publications

- **BriskSeed: Online History-Guided Reuse for Accelerating Dynamic Approximate Nearest Neighbor Search**. Hongru Gao, Shuhao Zhang, Haikun Liu, Xiaofei Liao, Hai Jin. ICDE 2027, Research First Round. Accepted.
- **BidKV: Utility-Guided Preemption Scheduling for KV-Pressure LLM Serving**. SC 2026. Corresponding Author.
- **SAGE: A Dataflow-Native Framework for Modular, Controllable, and Transparent LLM-Augmented Reasoning**. ICML 2026. Corresponding Author.
- **Neuromem: A Granular Decomposition of the Streaming Lifecycle in External Memory for LLMs**. ICML 2026. Corresponding Author.
- **CANDOR-Bench: Benchmarking In-Memory Continuous ANNS under Dynamic Open-World Streams**. SIGMOD 2026. Corresponding Author.
- **MorphStream: Adaptive Scheduling for Scalable Transactional Stream Processing on Multicores**. SIGMOD 2023 / TKDE 2025 line of work. Corresponding Author.
- **Towards Concurrent Stateful Stream Processing on Multicore Processors**. ICDE 2020. First Author.
- **BriskStream: Scaling Data Stream Processing on Shared-Memory Multicore Architectures**. SIGMOD 2019. First Author.

Honors and Service

- HUST Excellent Young Scholar; Huawei Donghu Young Scholar.
- VLDB Review Board Member; SC Program Committee Member; ICDE/TKDE Poster Track Chair.
- Program committee service for KDD, CIKM, DASFAA, SDM, and related venues.

Teaching and Open Materials

- Large Language Model Inference Systems and Practice, a practice-oriented graduate course connected to vLLM-HUST development.
- Graduate Paper Writing, public lecture materials for research writing and systems-paper methodology.