

# Towards Concurrent Stateful Stream Processing on Multicore Processors

**Shuhao Zhang**

co-author: Yingjun Wu (AWS), Fengzhang (Renmin University of China), Bingsheng He



Xtra Computing Group  
lead by Prof. Bingsheng He



# Data Stream Processing Systems

1. *BriskStream: Scaling Data Stream Processing on Shared-Memory Multicore Architectures, SIGMOD'19*
2. *StreamBox: Modern Stream Processing on a Multicore Machine, ATC'19*
3. *Grizzly: Efficient Stream Processing Through Adaptive Query Compilation, SIGMOD'20*

# Data Stream Processing Systems


- [Available] Fast stream processing on large-scale multicore architectures
  - E.g., BriskStream<sup>1</sup>, StreamBox<sup>2</sup>, Grizzly<sup>3</sup>
- [Inadequate] Support of *consistent stateful stream processing*

1. *BriskStream: Scaling Data Stream Processing on Shared-Memory Multicore Architectures*, SIGMOD'19
2. *StreamBox: Modern Stream Processing on a Multicore Machine*, ATC'19
3. *Grizzly: Efficient Stream Processing Through Adaptive Query Compilation*, SIGMOD'20

# Outline


- **Motivation**
- **State-of-the-art**
- **Our Approach**
- **Experimental Results**

# Motivation Example




Toll Processing:  
computes toll based on  
the road congestion  
status

# Motivation Example




Toll Processing:  
computes toll based on  
the road congestion  
status

# Motivation Example




Toll Processing:  
computes toll based on  
the road congestion  
status

# Motivation Example




Toll Processing:  
computes toll based on  
the road congestion  
status

# Motivation Example




- *Road congestion status is shared among different streaming operators*
- *Its consistency needs to be preserved*

# Motivation Example




- *Road congestion status is shared among different streaming operators*
- *Its consistency needs to be preserved*
  - *consistent stateful stream processing*

# The Current Common System Design




# The Current Common System Design



*Common designs:*

- a) *Pipelined processing with message passing*
- b) *On-demand data parallelism*

# The Current Common System Design




*Common designs:*

operator parallelism

- a) *Pipelined processing with message passing*
- b) *On-demand data parallelism*



# The Current Common System Design





*Common designs:*

- a) Pipelined processing with message passing*
- b) On-demand data parallelism*


# Key-based Stream Partition




# Key-based Stream Partition




# Key-based Stream Partition



# Key-based Stream Partition




# Key-based Stream Partition




Pay extra  
penalty if road1  
is empty!

# Key-based Stream Partition




Pay extra  
penalty if road1  
is empty!

# Key-based Stream Partition



Pay extra  
penalty if road1  
is empty!


# Key-based Stream Partition




Pay extra  
penalty if road1  
is empty!

*Can not efficiently handle general case*


# Lock-based State Sharing



# Lock-based State Sharing



# Lock-based State Sharing




*Lead to poor performance*

# Access Ordering


- A toll should based on exactly current road condition, otherwise..

# Access Ordering


- A toll should be based on **exactly current** road condition, otherwise..



# Access Ordering




# Access Ordering



*Common workaround:*

- **Buffer and sort**

# Access Ordering



*Common workaround:*


- **Buffer and sort**

*Lead to poor performance*


# Outline

- **Motivation**
- **State-of-the-art**
- **Our Approach**
- **Experimental Results**

# Consistent Stateful Stream Processing




# Consistent Stateful Stream Processing



## Employing transactional schematics.

- State transaction ( $txn_{ts}$ ):
  1.  $txn_{t1} \{Update\ Road1\ Cnt, update\ Roadx\dots\} @ 10:00;$
  2.  $txn_{t2} \{Read\ Road1\ Cnt, \dots\} @ 10:05$
- Correct schedule:  $txn_{t1} < txn_{t2} < \dots < txn_{tn}$


# Consistent Stateful Stream Processing



## Employing transactional schematics.

- State transaction ( $txn_{ts}$ ):
  1.  $txn_{t1} \{Update\ Road1\ Cnt, update\ Roadx\dots\} @ 10:00;$
  2.  $txn_{t2} \{Read\ Road1\ Cnt, \dots\} @ 10:05$
- Correct schedule:  $txn_{t1} < txn_{t2} < \dots < txn_{tn}$

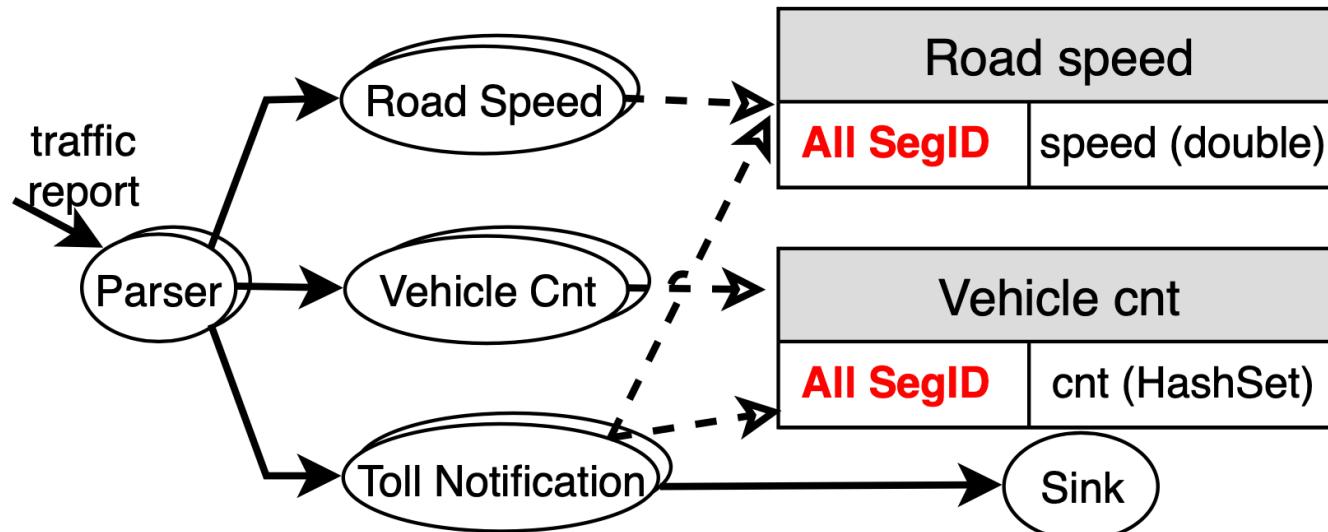
# Consistent Stateful Stream Processing



## Employing transactional schematics.

- State transaction ( $txn_{ts}$ ):
  1.  $txn_{t1} \{Update\ Road1\ Cnt, update\ Roadx\dots\} @ 10:00;$
  2.  $txn_{t2} \{Read\ Road1\ Cnt, \dots\} @ 10:05$
- Correct schedule:  $txn_{t1} < txn_{t2} < \dots < txn_{tn}$


# Consistent Stateful Stream Processing




## Employing transactional schematics.

- State transaction ( $txn_{ts}$ ):
  1.  $txn_{t1} \{Update\ Road1\ Cnt, update\ Roadx\dots\} @ 10:00;$
  2.  $txn_{t2} \{Read\ Road1\ Cnt, \dots\} @ 10:05$
- Correct schedule:  $txn_{t1} < txn_{t2} < \dots < txn_{tn}$

# Existing Solutions Revisited




# Existing Solutions Revisited



[1] S-Store: Streaming Meets Transaction Processing, VLDB'15

# Existing Solutions Revisited




What if we use existing design[1]?

- Severe lock contention

[1] S-Store: Streaming Meets Transaction Processing, VLDB'15

# Existing Solutions Revisited



What if we use existing design[1]?

- Severe lock contention

*A new solution for scaling concurrent state access is needed!*

[1] S-Store: Streaming Meets Transaction Processing, VLDB'15

# Outline

- **Motivation**
- **State-of-the-art**
- **Our Approach:**
  - TStream (An extension of BriskStream[2])
- **Experimental Results**

[2] BriskStream: Scaling Data Stream Processing on Shared-Memory Multicore Architectures, SIGMOD'19

# Design Overview

# Design Overview

- **Design 1) Dual-Mode Scheduling**
  - Exposing Parallelism
- **Design 2) Dynamic Restructuring Execution**
  - Exploiting Parallelism.

# Design Overview


- **Design 1) Dual-Mode Scheduling**
  - Exposing Parallelism
- **Design 2) Dynamic Restructuring Execution**
  - Exploiting Parallelism.

# Design Overview


- **Design 1) Dual-Mode Scheduling**
  - Exposing Parallelism
- **Design 2) Dynamic Restructuring Execution**
  - Exploiting Parallelism.

*Up to 4.8 times higher throughput with similar or even lower processing latency!*


## Design1) Dual-Mode Scheduling




# Design1) Dual-Mode Scheduling




## Design1) Dual-Mode Scheduling




# Design1) Dual-Mode Scheduling



## Design1) Dual-Mode Scheduling




## Design1) Dual-Mode Scheduling



- State access is often the bottleneck.
- It *unnecessarily* blocks all subsequent processes.


## Design1) Dual-Mode Scheduling



- State access is often the bottleneck.
- It *unnecessarily* blocks all subsequent processes.


*Let's postpone it.*

## Design 1) Dual-Mode Scheduling




(a)PreProcess

## Design 1) Dual-Mode Scheduling



## Design 1) Dual-Mode Scheduling




(a)PreProcess




(c)PostProcess


## Design 1) Dual-Mode Scheduling



(a)PreProcess




(c)PostProcess




(d)PreProcess

## Design 1) Dual-Mode Scheduling




## Design 1) Dual-Mode Scheduling



**Pros:** Enlarging inter-event parallelism opportunities.

**Cons:** Need to handle ``on-the-fly'' events.

## Design 1) Dual-Mode Scheduling



**Pros:** Enlarging inter-event parallelism opportunities.

**Cons:** Need to handle ``on-the-fly'' events.

# Design 2) Dynamic Restructuring Execution

## Design 2) Dynamic Restructuring Execution


**For each state transaction:**

Dynamically decompose and regroup operations to *avoid* conflict before actual parallel evaluation is applied.

## Design 2) Dynamic Restructuring Execution

**For each state transaction:**

Dynamically decompose and regroup operations to *avoid* conflict before actual parallel evaluation is applied.




**txn<sub>t1</sub>:** {**O<sub>1</sub>, O<sub>2</sub>**}  
**txn<sub>t2</sub>:** {**O<sub>3</sub>**}

## Design 2) Dynamic Restructuring Execution

**For each state transaction:**

Dynamically decompose and regroup operations to *avoid* conflict before actual parallel evaluation is applied.




$txn_{t1}: \{O_1, O_2\}$   
 $txn_{t2}: \{O_3\}$

$O_1$ : Read of Road 1  
 $O_2$ : Update to Road 2  
 $O_3$ : Read of Road 2

## Design 2) Dynamic Restructuring Execution

**For each state transaction:**

Dynamically decompose and regroup operations to *avoid* conflict before actual parallel evaluation is applied.




$txn_{t1}: \{O_1, O_2\}$   
 $txn_{t2}: \{O_3\}$

$O_1$ : Read of Road 1  
 $O_2$ : Update to Road 2  
 $O_3$ : Read of Road 2

## Design 2) Dynamic Restructuring Execution

**For each state transaction:**

Dynamically decompose and regroup operations to *avoid* conflict before actual parallel evaluation is applied.




$\text{txn}_{t1}: \{O_1, O_2\}$   
 $\text{txn}_2: \{O_3\}$

$O_1$ : Read of Road 1  
 $O_2$ : Update to Road 2  
 $O_3$ : Read of Road 2

## Design 2) Dynamic Restructuring Execution

**For each state transaction:**

Dynamically decompose and regroup operations to *avoid* conflict before actual parallel evaluation is applied.



$txn_{t1}: \{O_1, O_2\}$

$txn_2: \{O_3\}$

$O_1$  : Read of Road 1

$O_2$  : Update to Road 2

$O_3$  : Read of Road 2

# Outline

- **Motivation**
- **State-of-the-art**
- **Our Approach**
- **Experimental Results**

# Benchmark Workloads

# Benchmark Workloads

- **Four applications:**
  - Grep Sum (GS); Streaming Ledger (SL); Online Bidding (OB); Toll Processing (TP).
- **Diverse characteristics:**
  - Varying compute/state access ratio
  - Varying state transaction types: read-only, write-only
  - Varying data dependencies
  - Varying multi-partition transactions

# Benchmark Workloads

- **Four applications:**
  - Grep Sum (GS); Streaming Ledger (SL); Online Bidding (OB); Toll Processing (TP).
- **Diverse characteristics:**
  - Varying compute/state access ratio
  - Varying state transaction types: read-only, write-only
  - Varying data dependencies
  - Varying multi-partition transactions

# Benchmark Workloads

- **Four applications:**
  - Grep Sum (GS); Streaming Ledger (SL); Online Bidding (OB); Toll Processing (TP).
- **Diverse characteristics:**
  - Varying compute/state access ratio
  - Varying state transaction types: read-only, write-only
  - Varying data dependencies
  - Varying multi-partition transactions

# Benchmark Workloads

- **Four applications:**
  - Grep Sum (GS); Streaming Ledger (SL); Online Bidding (OB); Toll Processing (TP).
- **Diverse characteristics:**
  - Varying compute/state access ratio
  - Varying state transaction types: read-only, write-only
  - Varying data dependencies
  - Varying multi-partition transactions

# Benchmark Workloads


- **Four applications:**
  - Grep Sum (GS); Streaming Ledger (SL); Online Bidding (OB); Toll Processing (TP).
- **Diverse characteristics:**
  - Varying compute/state access ratio
  - Varying state transaction types: read-only, write-only
  - Varying data dependencies
  - Varying multi-partition transactions

# Benchmark Workloads


- **Four applications:**
  - Grep Sum (GS); Streaming Ledger (SL); Online Bidding (OB); Toll Processing (TP).
- **Diverse characteristics:**
  - Varying compute/state access ratio
  - Varying state transaction types: read-only, write-only
  - Varying data dependencies
  - Varying multi-partition transactions

Experiments conducted on a 4-socket Intel Xeon E7- 4820 server with 128 GB DRAM.


# Overall Performance Comparison




# Overall Performance Comparison




# Overall Performance Comparison




# Overall Performance Comparison




# Overall Performance Comparison



# Overall Performance Comparison




# Overall Performance Comparison




- (i) In general, TStream performs well at large core counts.
- (ii) TStream performs slightly better when data dependency is heavily presented (SL).
- (iii) Large room for further improvement.


# Overall Performance Comparison




(a) GS (50% read requests)



(b) SL




(c) OB



(d) TP

- (i) In general, TStream performs well at large core counts.
- (ii) TStream performs slightly better when data dependency is heavily presented (SL).
- (iii) Large room for further improvement.

# Overall Performance Comparison



- (i) In general, TStream performs well at large core counts.
- (ii) TStream performs slightly better when data dependency is heavily presented (SL).
- (iii) Large room for further improvement.

# Recap

## Recap

- [Key Designs] 1) dual-mode scheduling, 2) transaction restructuring
- [Results] TStream performs constantly better than prior solutions under varying workloads.

## Recap

- [Key Designs] 1) dual-mode scheduling, 2) transaction restructuring
- [Results] TStream performs constantly better than prior solutions under varying workloads.

<https://github.com/Xtra-Computing/briskstream/tree/TStream>

# Current Work



Next Generation IoT Data  
Management Platform  
(lead by Prof. Volker Markl)



visit us @  
[www.nebula.stream](http://www.nebula.stream)  
Twitter@nebulastream

# Current Work



Next Generation IoT Data  
Management Platform  
(lead by Prof. Volker Markl)



visit us @  
[www.nebula.stream](http://www.nebula.stream)  
Twitter@nebulastream

Thank you!  
shuhao.zhang@tu-berlin.de

# Competitor Schemes

## Competitor Schemes

- ① Ordering-lock-based approach (LOCK).
- ② Multi-versioning-based approach (MVLK).
- ③ Static-partition-based approach (PAT).
  - S-Store
- ④ No Consistency Guarantee (No-Lock).
  - Remove locks from LOCK scheme.
  - Represent the performance upper bound.

## Competitor Schemes

- ① Ordering-lock-based approach (LOCK).
- ② Multi-versioning-based approach (MVLK).
- ③ Static-partition-based approach (PAT).
  - S-Store
- ④ No Consistency Guarantee (No-Lock).
  - Remove locks from LOCK scheme.
  - Represent the performance upper bound.

## Competitor Schemes

- ① Ordering-lock-based approach (LOCK).
- ② Multi-versioning-based approach (MVLK).
- ③ Static-partition-based approach (PAT).
  - S-Store
- ④ No Consistency Guarantee (No-Lock).
  - Remove locks from LOCK scheme.
  - Represent the performance upper bound.

## Competitor Schemes

- ① Ordering-lock-based approach (LOCK).
- ② Multi-versioning-based approach (MVLK).
- ③ Static-partition-based approach (PAT).
  - S-Store
- ④ No Consistency Guarantee (No-Lock).
  - Remove locks from LOCK scheme.
  - Represent the performance upper bound.


## Competitor Schemes

- ① Ordering-lock-based approach (LOCK).
- ② Multi-versioning-based approach (MVLK).
- ③ Static-partition-based approach (PAT).
  - S-Store
- ④ No Consistency Guarantee (No-Lock).
  - Remove locks from LOCK scheme.
  - Represent the performance upper bound.


## Competitor Schemes

- ① Ordering-lock-based approach (LOCK).
- ② Multi-versioning-based approach (MVLK).
- ③ Static-partition-based approach (PAT).
  - S-Store
- ④ No Consistency Guarantee (No-Lock).
  - Remove locks from LOCK scheme.
  - Represent the performance upper bound.


## Runtime Breakdown




## Runtime Breakdown




# Runtime Breakdown




## Runtime Breakdown




## Runtime Breakdown




## Runtime Breakdown



## Runtime Breakdown




# Runtime Breakdown



- (a) No-Lock scheme spends more than 60% in others mainly due to index look up.
- (b) Synchronization overhead dominates in all prior solutions regardless of NUMA effect.
- (c) TStream involves high RMA overhead.

NUMA-aware optimization detailed in our paper.


# Runtime Breakdown



- (a) No-Lock scheme spends more than 60% in others mainly due to index look up.
- (b) Synchronization overhead dominates in all prior solutions regardless of NUMA effect.
- (c) TStream involves high RMA overhead.

NUMA-aware optimization detailed in our paper.

# Runtime Breakdown



- (a) No-Lock scheme spends more than 60% in others mainly due to index look up.
- (b) Synchronization overhead dominates in all prior solutions regardless of NUMA effect.
- (c) TStream involves high RMA overhead.

NUMA-aware optimization detailed in our paper.

## No free lunch

- **Currently, TStream has its drawback on efficiently handling transaction abort.**
  - E.g., one that leaves the average speed of a road to become negative.

## No free lunch

- **Currently, TStream has its drawback on efficiently handling transaction abort.**
  - E.g., one that leaves the average speed of a road to become negative.

*Stay tuned for future work*